

Volume-uniform electric-shell isolation, exact shared-link hopping, and a homological carrier in strong-coupling $SU(N)$ Hamiltonian lattice gauge theory

Master edition, with machine-checked provenance

Alexander Smith
Independent Researcher

30 August 2026

Abstract

We study the Kogut–Susskind Hamiltonian for pure $SU(N)$ gauge theory on a periodic cubic spatial lattice at strong coupling. In the normalised charge-odd fundamental-plaquette sector, an explicit order-two Weingarten calculation gives the exact coefficient of hopping between adjacent plaquettes that share one link,

$$t_N = \frac{2N(N^2 - 4)}{(N^2 - 1)(2N^2 - 1)(4N^2 - 9)} > 0, \quad N \geq 3.$$

The oriented cellular boundary $B = \partial_2^\dagger$ fixes the corresponding local geometry. A Casimir and short-cycle argument proves the volume-uniform electric window $\text{spec}(H_E|_{\text{trivial flux}}) \cap [0, \frac{5}{2}C_F) = \{0, 2C_F\}$, with the $2C_F$ eigenspace exactly the oriented one-plaquette shell and an external margin at least $C_F/2$. A link-centre classification then proves that the four shared-link fusion channels exhaust every second-order off-diagonal process. Consequently the global order- u^2 symbol is $t_N u^2 B(k)B(k)^\dagger$, with no retained-shell or process-completeness hypothesis. Its spectrum is $\{0, t_N u^2 q, t_N u^2 q\}$ up to a common scalar, with $q = 4 \sum_j \sin^2(k_j/2)$, and its flat carrier $\ker \partial_2$ has dimension $L^3 + 2$. A normalised Fourier sum of radius-one cube boundaries is the exact carrier at every nonzero allowed momentum; on nested volumes its order- u^2 splitting is L -independent, although one local cube seed has only $O(L^{-3})$ weight at a fixed momentum.

A supplied $SU(3)$ two-cube calculation — an exact-Casimir $\mathcal{B} = 6$ basis of 1,590,462 states, reached by path enumeration without ever assembling the magnetic matrix — reports the same coefficient channel by channel rather than only in the sum, with each of the six link-resolved channels separately proportional to the same incidence Gram on all 56 adjacent pairs, and localises the $\mathcal{B}$ -cutoff dependence to one symmetry orbit; those finite-precision contractions are corroboration, not an independent derivation. By arithmetic on the weight ledger alone, a cutoff keeping only the singlet and $\Lambda^2 F$ routes projects the hopping to $-1/12$ — opposite in sign to t_3 — and omits exactly 14/153, the two summing to t_3 . A later detached replay on the same open prism yields nested radius-two finite- u Ritz spectra and the leading incremental order- u^4 and order- u^5 effects of entering Q_2 and restoring its retained magnetic block; these are finite-volume nested-model differences, not complete fourth- or fifth-order coefficients. Exact integer geometry also gives the closed-cube shell $A_1^{-} \oplus T_1^{+-} \oplus E^{--}$ with Gram eigenvalues 0, 4, 6.

For comparison we solve the spectrum of a defined unsigned-incidence operator in closed form, the cubic $\mu(\mu-p)^2 = 4\hat{a}_1\hat{a}_2\hat{a}_3$ with $p+q = 12$, so one zone function runs both sectors. Identifying that operator with the complete charge-even linked Hamiltonian requires a further derivation,

because the vacuum-mediated route must be accounted for. The organized upstream corpus records a cold exact replay of the retained $SU(3)$ third-order chain, although that microscopic replay is not bundled here. Fourth-order physical-kernel diagnostics remain conditional; in particular, the retained checkpoints provably do not determine the off-axis coefficient, and that non-identifiability is exhibited by an explicit second witness.

The exact coefficient results are finite-volume and finite-order. In addition, for $SU(3)$ a published spectral-localisation theorem gives a volume-uniform, all-orders Riesz island for the complete neutral charge-odd shell at sufficiently small local coupling; the threshold is existential and the result remains at fixed lattice spacing. No continuum glueball or Yang–Mills mass-gap claim is made.

1 Introduction

Hamiltonian strong-coupling expansions organise lattice gauge dynamics around the electric vacuum and its flux excitations [3, 4]. In three spatial dimensions a fundamental plaquette excitation carries one of three face orientations, and at second order neighbouring faces communicate through a shared link, so the projected problem resembles a three-orbital tight-binding model. The signs of those matrix elements are fixed by the orientation of the plaquette boundary and by charge conjugation, and the resulting interference is the central object here.

Flat bands generated by incidence or Gram matrices have a long history: local topology and signed or line-graph constructions [2, 12, 16, 20, 22], compact localised states and their failure to span singular bands on a torus [19, 21], and supersymmetric incidence factorisations [25] are all established. Closely related closed-membrane structures also occur in two-form spin liquids [26]. A recent classification of exactly flat bands protected by local graph topology [30] reaches extensive face-graph degeneracy from an *unoriented* face–vertex incidence, where rank–nullity alone gives $|F| - |V|$ flat bands. That is a different mechanism from the one here, and the difference is measurable: on the signed face–edge operator of Section 3 the corresponding index is exactly 0 and bounds nothing, while both kernels equal $L^3 + 2$ — homology, not counting. It is recorded as a distinction, not an identification. We do not claim that incidence flatness, cellular compact localised states, or homological counting are new in isolation.

The new dynamical content is local and exact. Starting from non-Abelian $SU(N)$ Hamiltonian lattice gauge theory, the shared-link matrix element in the charge-odd fundamental-plaquette sector factorises into a colour coefficient and an oriented edge–face Gram entry. Representation theory fixes the former and the cubical chain complex fixes the latter,

$$\begin{aligned} \text{colour dynamics} &\longrightarrow t_N, \\ \text{cellular geometry} &\longrightarrow BB^\dagger. \end{aligned}$$

What is proved, and what remains conditional

The principal second-order statements are self-contained here; where a higher-order or charge-even operator’s identification with a physical sector is conditional, that condition is stated separately. Earlier editions carried three load-bearing prose premises. Two are now theorems, and this edition proves more than the premises asked for: the free electric window of Theorem 1 and the process classification of Lemma 16. The third, the charge-even identification, remains a hypothesis and is argued in Section 11; the coverage appendix (Appendix G) makes its unchecked status visible by omission.

1. For every $N \geq 3$ the shared-link channel weights are d_ρ/N^2 , derived from the order-two

Weingarten values (Theorem 12). This was previously an isotropy *hypothesis*, and everything in Section 4 rests on it.

2. For every $N \geq 3$ and $L \geq 3$, the trivial-flux electric spectrum below $\frac{5}{2}C_F$ is exactly $\{0, 2C_F\}$; the charge-odd part of the $2C_F$ eigenspace is exactly the $3L^3$ -dimensional one-plaquette shell and its external free margin is at least $C_F/2$ (Theorem 1). Every off-diagonal second-order intermediate is then either a vanishing same-face route or one of the four adjacent shared-link channels (Lemma 16). Theorem 17 gives their explicit projector cross matrix elements and coefficient t_N , so the global symbol $t_N u^2 B B^\dagger$ follows without a retained-shell or process-completeness hypothesis (Corollary 18).
3. The Bloch spectrum of $B B^\dagger$ is $\{0, q, q\}$ (Theorem 5), and on the periodic L^3 complex the zero-mode count is $L^3 + 2$ (Theorem 6). The two are the same statement, by a sum rule over the momentum grid (Theorem 8). At fixed nonzero nested momentum the carrier and its splitting are volume-independent, with the exact locality–momentum tradeoff given in Proposition 10.
4. The truncation that keeps only the singlet and $\Lambda^2 F$ routes projects the hopping to $w_{\Lambda^2 F} - w_1 = -1/12$, and what it omits is $w_{\text{Sym}^2 F} - w_{\text{Adj}} = 14/153$; the two sum to t_3 (Corollary 23). This is arithmetic on (32) and needs no input beyond Theorem 12.
5. The defined unsigned-incidence comparison operator has the cubic $\mu(\mu - p)^2 = 4\hat{a}_1\hat{a}_2\hat{a}_3$ with $p + q = 12$ (Theorem 31); its range is exactly $[-4, 12]$, the top attained only at Γ and the floor on the three planes $k_j = \pi$, triple only at R (Proposition 32). Conditional on identifying this comparison operator with the complete linked charge-even symbol, its Γ -point splitting is $12t_{r,+}$ (Proposition 34). That identification is not derived here; the vacuum-mediated all-pairs term must first be reorganised or cancelled.
6. On the closed cube surface the charge-odd shell is $A_1^{--} \oplus T_1^{+-} \oplus E^{--}$, with G -eigenvalues 0, 4, 6; the labels follow from the rotation action rather than being assigned, the A_1^{--} level is the fundamental class of S^2 , and it sits at the scalar for every hopping coefficient (Proposition 26). This is integer linear algebra and takes no input from any delivery.
7. The retained Γ /axis corpus cannot identify the disputed fourth-order off-axis coefficient (Theorem 40), and an explicit second witness exhibits the non-identifiability by construction (Proposition 42).

Four families of statements depend on computational artifacts. The radius-two parents and detached verifier are bundled with this edition; the earlier raw class contractions, symbolic history certificates and microscopic rank-sweep artifacts are not. Repeated assertions in synthesis documents are not treated as independent reproductions. The organized upstream authority corpus records a cold exact 251/251 replay for the retained third-order operator chain; that is stronger than an unchecked delivery, but it is still cited here as external computational evidence because this release does not replay the raw chain.

8. A supplied two-cube delivery reports six link-irrep channel coefficients at $N = 3$, as rational reconstructions of finite-precision contractions. Conditional on those six numbers, they are the four fusion-channel weights individually — channel by channel and not merely in the sum (Reported result 21) — and the $\mathcal{B} = 4$ truncation of that space retains exactly the adjacent shared-link channels the published link cutoff of item 4 retains, so the two land on the same projected hopping — the same channels, not the same scheme. Conditional on the delivered

diagonals, the symmetry analysis of Proposition 25 applies to them: the orbit structure itself is exact and takes no delivered input. The $\mathcal{B} = 6$ two-cube build was not reproduced here.

9. A later hash-sealed release on the same open prism reports nested charge-odd and charge-even radius-two Ritz operators. Its detached verifier reconstructs the spectra and the leading $K_{2,\star} - K_1$ and $K_{2,\text{full}} - K_{2,\star}$ coefficients from the sealed parity-sector parents. This establishes numerical consistency relative to those inputs, not convergence in action radius or a physical glueball gap (Reported result 28).
10. For $SU(3)$, the frozen upstream corpus records that all candidate third-order lifter classes vanish and gives the scalar and shared-link coefficients of Section 7. The census is three lifter classes evaluated over 32 five-trace numerator patterns each, 96 contractions in total, within a retained chain that passed 251/251 cold exact gates upstream. This compact verifier checks the resulting rational assembly but not those microscopic contractions. If the census is exhaustive and the linked operator ledger complete, the effective Hamiltonian stays in the incidence algebra through order u^3 .
11. Supplied ledgers report the all-rank cube formula $\alpha_N = 640/[N(N^2 - 1)^3]$ and its equality with the complete axial coefficient for $N = 3, \dots, 12$. Conditional on those calculations, the $SU(3)$ carrier acquires nonzero axial dispersion at order u^4 . The complete off-axis operator remains open — and by Theorem 40 it cannot be closed by any Γ or axial datum.

These coefficient statements concern the effective Hamiltonian projected to the degenerate one-plaquette manifold. Virtual multi-plaquette states are retained in the resolvents, but they do not by themselves solve the external spectral problem. At $k = 0$ the three incidence branches touch and the finite-volume separation collapses as L^{-2} . Theorem 2 later proves all-orders persistence of the complete three-component cluster in an existential small- $|u|$ domain; it does not prove convergence of the displayed coefficient series, persistence of one internal sheet, identification with the lightest physical glueball, or a continuum mass gap.

2 Hamiltonian and projected sector

Let $\Lambda_L = (\mathbb{Z}/L\mathbb{Z})^3$, $L \geq 3$, be a cubic spatial lattice with periodic boundary conditions. In units of the electric coupling the Kogut–Susskind Hamiltonian is

$$H_\beta = \frac{1}{2} \sum_{\ell} C_2(\ell) + \beta_N \sum_p \left(1 - \frac{1}{N} \text{Re Tr } U_p\right). \quad (1)$$

Dropping the additive constant,

$$H = H_E - uV, \quad V = \sum_p (\chi_p + \bar{\chi}_p), \quad u = \frac{\beta_N}{2N}, \quad (2)$$

with $\chi_p = \text{Tr } U_p$. For $SU(3)$, $u = \beta_3/6$. The definition $u = \beta_N/(2N)$, not the $SU(3)$ -specific equality, is the all-rank convention; the archived line $Y = 2\beta/3 = 4u$ is a definition-label erratum and the printed coefficients are already in u . Reading it as a coordinate would demand dividing the order- r coefficient by 4^r , which misses the printed tables by exactly 16 at u^2 and 64 at u^3 .

All quoted excitation energies use the linked, vacuum-subtracted operator $H_{\text{exc}}(u) = H(u) - E_{\text{vac}}(u)I$. This changes only the scalar ledger; incidence hopping and carrier shape are unaffected.

The electric vacuum $|0\rangle$ carries the trivial representation on every link. For $N \geq 3$ the normalised one-plaquette states of definite charge conjugation are

$$|p, \pm\rangle = \frac{\chi_p \pm \bar{\chi}_p}{\sqrt{2}} |0\rangle, \quad (3)$$

of unit norm by character orthogonality. Each of the four boundary links carries the fundamental or antifundamental representation, so

$$E_{F,N} = 4\left(\frac{1}{2}C_F\right) = 2C_F = \frac{N^2 - 1}{N}, \quad C_F = \frac{N^2 - 1}{2N}. \quad (4)$$

This external energy is the input to the resolvent denominators of Section 4, where it is checked together with them. For $SU(2)$ complex conjugation is an inner gauge transformation and the charge-odd state vanishes; this is why the construction begins at $N = 3$. The same fact appears again in Section 4: $\Lambda^2 F$ is the singlet exactly at $N = 2$, which is where t_N vanishes. The analytic continuation of the unsigned-channel coefficient is nonzero there, which locates the $N^2 - 4$ zero in the channel difference rather than in the incidence. Because the charge-conjugation split degenerates for $SU(2)$, this continuation is not asserted to be an $SU(2)$ charge-even spectrum.

For each coordinate direction choose a dual noncontractible cut Σ_j , and fix one orientation for every geometric edge. For $z \in \mathbb{Z}_N$, let $Z_j(z)$ act by $U_e \mapsto z^{\varepsilon(e, \Sigma_j)} U_e$, where the signed intersection number $\varepsilon(e, \Sigma_j) \in \{-1, 0, +1\}$. These unitaries commute with H_E , while charge conjugation obeys $CZ_j(z)C^{-1} = Z_j(z^{-1})$. The simultaneous $+1$ eigenspace of every $Z_j(z)$ is therefore charge-conjugation invariant; call it the *trivial electric one-form sector*. A loop of N -ality k and winding vector w has eigenvalue z^{kw_j} , so contractible plaquettes lie in this sector while a once-winding fundamental loop does not.

Theorem 1 (Uniform first electric spectral window). *Let $N \geq 3$ and $L \geq 3$. Let the pure-gauge physical Hilbert space be $\mathcal{H}_{\text{phys}} = L^2(SU(N)^{E(\Lambda_L)})^{SU(N)^{V(\Lambda_L)}}$, and define the trivial electric one-form sector and its charge-odd part by*

$$\mathcal{H}_{\text{phys}}^{(0)} = \{\Psi \in \mathcal{H}_{\text{phys}} : Z_j(z)\Psi = \Psi \ \forall j, z\}, \quad \mathcal{H}_{\text{phys}}^{(0,-)} = \{\Psi \in \mathcal{H}_{\text{phys}}^{(0)} : C\Psi = -\Psi\}.$$

Normalise the electric Hamiltonian on the Peter–Weyl spin-network basis by $H_E\Psi_{\{R_e\}, \{\iota_v\}} = \frac{1}{2} \sum_e C_2(R_e)\Psi_{\{R_e\}, \{\iota_v\}}$. Then, with $E_{F,N} = 2C_F$,

$$\text{spec}(H_E|_{\mathcal{H}_{\text{phys}}^{(0)}}) \cap \left[0, \frac{5}{2}C_F\right) = \{0, 2C_F\}. \quad (5)$$

The zero eigenspace is the vacuum line, while the $2C_F$ eigenspace is the direct sum, over elementary plaquettes, of their two oriented fundamental Wilson-loop lines. Consequently

$$\text{Ran}\left(\mathbf{1}_{\{E_{F,N}\}}(H_E)|_{\mathcal{H}_{\text{phys}}^{(0,-)}}\right) = \text{span}\{|p, -\rangle : p \in F_2(\Lambda_L)\}, \quad (6)$$

and this charge-odd eigenspace has dimension $3L^3$. In particular the full electric shell is separated from the rest of the trivial-flux spectrum by

$$\text{dist}\left(2C_F, \text{spec}(H_E|_{\mathcal{H}_{\text{phys}}^{(0)}}) \setminus \{2C_F\}\right) \geq \frac{1}{2}C_F, \quad (7)$$

uniformly in L .

Proof. Normalise the invariant form so every root has squared length 2, and write a nonzero dominant highest weight as $\lambda = \sum_{i=1}^{N-1} a_i \omega_i$. If $a_i > 0$, put $\lambda = \omega_i + \mu$ with μ dominant. Positivity of the inverse A_{N-1} Cartan matrix gives $C_2(\lambda) - C_2(\omega_i) = \frac{1}{2}((\mu, \mu) + 2(\omega_i + \rho, \mu)) \geq 0$, with equality only for $\mu = 0$. Moreover

$$C_2(\omega_i) = \frac{i(N-i)(N+1)}{2N}, \quad C_2(\omega_i) - C_F = \frac{N+1}{2N}(i-1)(N-i-1) \geq 0.$$

Thus every nontrivial irrep satisfies $C_2(R) \geq C_F$, with equality only for F and $\bar{F}$.

We need one further all-rank shelf. If a dominant weight λ is neither ω_1 nor ω_{N-1} , then one of three cases occurs: an interior Dynkin coefficient is nonzero; an endpoint coefficient is at least two; or both endpoint coefficients are nonzero. Monotonicity of the preceding Casimir difference therefore bounds $C_2(\lambda)$ below by, respectively, $C_2(\omega_2)$ (for $N \geq 4$), $C_2(2\omega_1)$ or its conjugate, or $C_2(\omega_1 + \omega_{N-1})$. The three comparisons with the required shelf are

$$C_2(\omega_2) - \frac{5}{4}C_F = \frac{(N+1)(3N-11)}{8N} > 0 \quad (N \geq 4), \quad (8)$$

$$C_2(2\omega_1) - \frac{5}{4}C_F = \frac{(N-1)(3N+11)}{8N} > 0, \quad (9)$$

$$C_2(\omega_1 + \omega_{N-1}) - \frac{5}{4}C_F = \frac{3N^2+5}{8N} > 0. \quad (10)$$

For $N = 3$ there is no interior coefficient, so the final two cases are exhaustive. Hence every irrep other than $\mathbf{1}, F, \bar{F}$ satisfies

$$C_2(R) > \frac{5}{4}C_F. \quad (11)$$

We also need the centre-neutral refinement. Zero N -ality is equivalent to λ lying in the A_{N-1} root lattice. Write $\lambda = \sum_i m_i \alpha_i$, $m_i \in \mathbb{Z}$, and $a_i = (\lambda, \alpha_i) \geq 0$. Since $m = A^{-1}a$ and every entry of A^{-1} is strictly positive, a nonzero dominant λ has $m_i \geq 1$ for every i . Hence $(\lambda, \rho) = \sum_i m_i \geq N-1$. Every nonzero root-lattice vector has squared norm at least 2, so

$$C_2(\lambda) = \frac{1}{2}((\lambda, \lambda) + 2(\lambda, \rho)) \geq N. \quad (12)$$

Consider a non-vacuum spin-network component with energy $E < \frac{5}{2}C_F$. If S is the set of links carrying nontrivial irreps, then $\sum_{e \in S} C_2(R_e) = 2E < 5C_F$, so $|S| \leq 4$. Gauge invariance forbids support degree one, because a single nontrivial irrep has no invariant vector. The torus graph is simple for $L \geq 3$; therefore a nonempty connected support has at least a three-cycle. Two components would require at least six links, and a connected simple graph of minimum degree two with at most four edges is a single C_3 or C_4 .

A three-cycle occurs only for $L = 3$ and is a straight loop winding once around one coordinate direction. Indeed, lift its three unit steps to $\mathbb{Z}^3$: their sum is Lw , and the length and parity constraints force $L = 3$, $w = \pm e_j$. Degree-two intertwiners transport one irrep R around the loop. Trivial one-form flux forces its N -ality to vanish, and (12) gives $H_E = \frac{3}{2}C_2(R) \geq \frac{3}{2}N > \frac{5(N^2-1)}{4N} = \frac{5}{2}C_F$, a contradiction. The support is therefore a four-cycle. Degree-two intertwiners transport one irrep R (with the dual on oppositely oriented links) around it, so its energy is $2C_2(R)$. Lift its four unit steps to $\mathbb{Z}^3$. For $L \geq 5$ the winding is zero; at $L = 3$ a nonzero winding violates step parity; and at $L = 4$ the only nonzero possibility is four collinear steps, a once-winding loop. Trivial flux then forces R to have zero N -ality, so (12) gives $H_E = 2C_2(R) \geq 2N > \frac{5}{2}C_F$, again a contradiction.

With zero winding, simplicity forces, up to cyclic permutation, reversal, coordinate relabelling, and independent sign choices, the cyclic step sequence $\varepsilon_i e_i, \varepsilon_j e_j, -\varepsilon_i e_i, -\varepsilon_j e_j$, with $i \neq j$: the

boundary of one elementary plaquette. Equation (11) now forces $R \in \{F, \bar{F}\}$, and its energy is $2C_F$. This proves (5); the empty support gives the unique vacuum.

At each degree-two vertex Schur's lemma makes the intertwiner unique. The two orientations are $\chi_p|0\rangle$ and $\bar{\chi}_p|0\rangle$, exchanged by charge conjugation, so their negative eigenspace is the line $|p, -\rangle$. Conversely every such state is physical, contractible, and has energy $4(C_F/2)$. Peter–Weyl orthogonality separates distinct plaquettes, of which there are $3L^3$. Finally, the vacuum lies $2C_F$ below the shell and (5) places every higher eigenvalue at least $C_F/2$ above it, proving (7). $\square$

The preceding theorem is enough to identify the physical content of a small-coupling spectral patch. The volume-uniform persistence of that patch is not obtained by bounding the norm of the extensive perturbation. Instead we use the local spectral enclosure of Yarotsky [1].

Theorem 2 (*$SU(3)$ all-orders finite-volume Riesz island*). *For periodic L^3 tori with $L \geq 3$, let $\tilde{H}_L(u) = H_L(u) - E_{0,L}(u)I$ and $\lambda = 27c_2|u|$, where $c_1, c_2 > 0$ are the range-dependent constants in the spectral localisation theorem of [1]. If*

$$27|u| < c_1, \quad \lambda < \frac{1}{33}, \quad (13)$$

then the part of $\text{spec } \tilde{H}_L(u)$ descending from the free energy $E_F = 8/3$ is separated from every other free-energy patch by

$$\delta_-(u) \geq \frac{1 - 15\lambda}{3}, \quad \delta_+(u) \geq \frac{1 - 33\lambda}{6}, \quad (14)$$

uniformly in L . The circle

$$\Gamma_F : |z - 8/3| = \rho, \quad \rho = \frac{1 - \lambda}{12}, \quad (15)$$

encloses that patch and obeys

$$\sup_{z \in \Gamma_F} \|(\tilde{H}_L(u) - z)^{-1}\| \leq \frac{12}{1 - 33\lambda}. \quad (16)$$

Restricted to $\mathcal{H}_{\text{phys}}^{(0,-)}$, its Riesz projection has rank exactly $3L^3$.

Proof. Group the three positively oriented outgoing links at each vertex into one cell and rescale the Hamiltonian by $3/2$. The smallest nonzero one-link electric energy is $2/3$, so the rescaled one-cell Hamiltonian has gap one. Assign the three positively oriented plaquettes based at a vertex to that cell. Since $\|\chi_p + \bar{\chi}_p\| \leq 6$, the norm of each rescaled cell interaction is bounded by $\frac{3}{2}|u| \cdot 3 \cdot 6 = 27|u|$.

The cell spaces may be infinite dimensional and the onsite Hamiltonians may be unbounded in the theorem of [1]. Its linked-set proof uses only finite range, the local norm, and an exponential bound on the number of connected sets through a marked cell. The same bound holds on every periodic finite quotient because its interaction graph has the same uniformly bounded degree. Thus the constants are independent of L and, translated back to the present energy units,

$$\text{spec } \tilde{H}_L(u) \subset \bigcup_{a \in \text{spec } H_{E,L}} [a(1 - \lambda), a(1 + \lambda)]. \quad (17)$$

For an $SU(3)$ irrep (p, q) , its one-link electric energy in units of $1/6$ has numerator $p^2 + q^2 + pq + 3p + 3q$. The nonzero one-link numerators below 18 are 4, 9, 10, 16, so their additive semigroup has consecutive attainable values 14, 16, 17 around the target 16. Consequently the lower and upper separations between the target interval and its nearest competitors are exactly the right-hand sides of (14); the upper one is positive precisely under the second condition in (13). The target

interval has half-width $8\lambda/3$. Direct subtraction from (15) gives the minimum contour clearance $(1 - 33\lambda)/12$, proving (16).

Gauge transformations, charge conjugation and the electric one-form symmetries commute with $H_L(u)$ and hence with the contour projection. The same contour works along $t \mapsto tu$, $0 \leq t \leq 1$. Finite-volume compact resolvent and bounded analytic perturbation make the restricted projection norm-continuous, so its finite rank is constant. At $u = 0$ Theorem 1 identifies that rank as $3L^3$. $\square$

Remark 3 (Exact boundary of the all-orders result). Theorem 2 transports the complete three-component shell, not one internally isolated band. The constants c_1, c_2 are not numerical, the result is confined to fixed-spacing strong coupling, and the separate one-particle theorem of [1] is not invoked: its nondegenerate one-site, nonresonance hypothesis fails for this composite, degenerate plaquette shell. The theorem therefore supplies an interacting finite-volume Riesz cluster, not an infinite-volume particle or Wilson-source residue.

Remark 4 (The two small-torus coincidences). At $L = 4$ a once-winding fundamental loop has energy $2C_F$ for every $N \geq 3$; trivial one-form flux is therefore load-bearing at every rank, not only at $SU(4)$. The genuinely rank-four coincidence occurs instead at $L = 3$: an $SU(4)$ loop in $\Lambda^2 F$ has $C_2 = 5/2$ and energy $3C_2/2 = 15/4 = 2C_F$, but its N -ality is two, so the same flux condition excludes it (and its character is charge-even). These are volume effects, not a continuum statement.

Set P_- equal to the projector in (6) and $Q_- = 1_- - P_-$, where 1_- is the identity on $\mathcal{H}_{\text{phys}}^{(0,-)}$. Both H_E and V preserve this space: H_E commutes with C and every $Z_j(z)$, while every contractible plaquette term in V is charge-even and centre-neutral. No degenerate $E_{F,N}$ direction is omitted, and degenerate perturbation theory gives

$$H_{\text{eff},-} = P_- H_E P_- + u H_1 + u^2 H_2 + u^3 H_3 + \dots \quad (18)$$

with every H_r linked, vacuum-subtracted, and carrying the appropriate Q -space resolvents.

3 Oriented incidence and the exact carrier

The periodic cubical complex has the cellular chain sequence $C_3 \xrightarrow{\partial_3} C_2 \xrightarrow{\partial_2} C_1$ with $\partial_2 \partial_3 = 0$. One-plaquette amplitudes live in C_2 and their signed leakage into links is $\partial_2 c$: a cellular closed-surface condition, not the microscopic vertex Gauss law, which the Wilson-loop basis already satisfies through $\partial_1 \partial_2 = 0$.

Writing $d_j = e^{ik_j} - 1$ and $q(k) = \sum_j |d_j|^2 = 4 \sum_j \sin^2(k_j/2)$, the face-to-link Bloch incidence matrix is

$$B(k) = \partial_2(k)^\dagger = \begin{pmatrix} d_2 & -d_1 & 0 \\ d_3 & 0 & -d_1 \\ 0 & d_3 & -d_2 \end{pmatrix}. \quad (19)$$

Theorem 5 (Incidence factorisation). *With $w(k) = (\bar{d}_3, -\bar{d}_2, \bar{d}_1)^\top$,*

$$BB^\dagger = qI - ww^\dagger, \quad B^\dagger w = 0, \quad \|w\|^2 = q,$$

and therefore $\text{spec}(BB^\dagger) = \{0, q, q\}$. With $S = BB^\dagger - 4I$ the signed shared-link adjacency, $\text{spec } S(k) = \{-4, q - 4, q - 4\}$.

Proof. Multiplying (19) by its adjoint gives the first identity entry by entry. Since $\|w\|^2 = q$, the rank-one operator $qI - ww^\dagger$ annihilates w and acts as qI on its orthogonal complement. Subtracting $4I$ gives the adjacency spectrum. $\square$

For $k \neq 0$ the carrier is one-dimensional and represented by $w(k)$. At Γ , $B = 0$ and all three branches meet; the normalised Bloch eigenvector has no direction-independent limit there, which is the familiar singular-flat-band mechanism behind the incompleteness of a purely compact localised basis.

Theorem 6 (Finite-torus carrier). *For the periodic cubic cellulation of T^3 , $\ker \partial_2 = \text{im } \partial_3 \oplus H_2$ with $\dim \text{im } \partial_3 = L^3 - 1$ and $\dim H_2 = 3$, so $\dim \ker \partial_2 = L^3 + 2$.*

Proof. There are L^3 three-cells and no four-cells, and $b_3(T^3) = 1$, so $\text{rank } \partial_3 = L^3 - 1$ by rank-nullity. Discrete Hodge decomposition gives $\ker \partial_2 = B_2 \oplus H_2$ with $\dim H_2 = b_2(T^3) = 3$. $\square$




Figure 1: The carrier is a cellular cycle space. The left panel shows three visible terms in the outward-oriented boundary of one cube; the complete six-face formula is given in Appendix E. Cube boundaries span $\text{im } \partial_3$, while three noncontractible wrapping classes complete $\ker \partial_2$ on the three-torus.

The count is verified over $\mathbb{Z}$ rather than re-derived from the formula: the complex is built from the boundary formulas of Appendix E; the exhibited cycles — every elementary cube boundary together with three wrapping sheets — bound the kernel from below, and rank over $\mathbb{F}_p$ bounds it from above, since rank can only drop under reduction. The two bounds meet at $L = 1, \dots, 5$.

Remark 7. $L^3 + 2$ is the nullity of the down boundary Laplacian, not a Betti number; the Betti contribution is the final 3. The restriction $L \geq 3$ belongs to the twelve-neighbour Bloch adjacency, where $x + e$ and $x - e$ coincide at $L = 2$. The chain-level count has no such restriction and holds at $L = 1, 2$ as well.

The following sum rule joins Theorems 5 and 6, which are otherwise two unrelated statements: a 3×3 spectrum and an $(L^3 + 2)$ -dimensional space.

Theorem 8 (Bloch-chain bridge). *On the periodic L^3 complex,*

$$\sum_k (3 - \text{rank } B(k)) = L^3 + 2,$$

the sum running over the L^3 allowed momenta, with nullity profile 3 at Γ exactly once and 1 at each of the other $L^3 - 1$ momenta.

Proof. $B(0) = 0$ contributes 3. For $k \neq 0$ some $d_j \neq 0$, and Theorem 5 gives $\text{rank } B(k) = 2$, contributing 1 each. Summing gives $3 + (L^3 - 1) = L^3 + 2$. $\square$

Remark 9. Earlier editions of this paper read the agreement as a coincidence worth checking — “the 3 is the triple degeneracy of $B(0) = 0$, not $b_2(T^3)$ ”. That was wrong, and the correction is the better statement. Under the discrete Fourier transform ∂_2 block-diagonalises and $\ker \partial_2 = \bigoplus_k \ker \partial_2(k)$; at $k \neq 0$ the kernel block is one-dimensional and is exactly $\text{im } \partial_3(k)$, while $\partial_3(0) = 0$. So all of $\text{im } \partial_3$ sits at nonzero momenta and the quotient is supported entirely at Γ : the Γ block is H_2 , and the 3 is $b_2(T^3)$. The two routes do not agree by accident; they are the same decomposition read in two bases.

The same block decomposition says exactly what a local seed costs at sharp momentum.

Proposition 10 (Fixed-momentum carrier and the locality tradeoff). *Let $b_x = \partial_3[x; 1, 2, 3]$ be the six-face boundary of the elementary cube at x , and use the unitary discrete Fourier transform $\hat{b}_L(k) = L^{-3/2} \sum_{x \in \Lambda_L} e^{-ik \cdot x} b_x$. For every allowed $k \neq 0$, $\hat{b}_L(k)$ has norm $\sqrt{q(k)}$ and, up to the Fourier phase convention, is the vector $w(k)$ of Theorem 5. Hence $\varphi_L(k) = \hat{b}_L(k)/\sqrt{q(k)}$ is exactly the unique unit carrier vector in that momentum fibre.*

If $k_0 = 2\pi(r_1, r_2, r_3)/L_0 \neq 0$, then k_0 is an allowed momentum on every nested volume $L = mL_0$, and the order- u^2 internal separation from each of the two incidence branches is

$$t_N u^2 q(k_0), \quad (20)$$

independent of m . More generally it is bounded below by $t_N u^2 q_$ on any set of allowed momenta with $q(k) \geq q_* > 0$.*

This exact sharp-momentum state is not local: it is the Fourier sum of all cube boundaries. Conversely, for the normalised radius-one seed $b_x/\sqrt{6}$,

$$|\langle \varphi_L(k), b_x/\sqrt{6} \rangle| = \sqrt{\frac{q(k)}{6L^3}}, \quad (21)$$

and $\hat{b}_L(0) = 0$.

Proof. Fourier transforming the six terms of the cubical boundary gives, in the $(12, 13, 23)$ face basis, $(d_3, -d_2, d_1)^T$, with conjugates if the opposite Fourier sign is chosen. Its squared norm is $q(k)$ and $\partial_2 \partial_3 = 0$ puts it in $\ker B(k)^\dagger$; uniqueness for $k \neq 0$ follows from Theorem 5. The nested-volume statement is the integer identity $k_0 = 2\pi(mr)/(mL_0)$, and (20) follows from Corollary 18. Inverting the unitary Fourier transform gives $b_x = L^{-3/2} \sum_k e^{ik \cdot x} \sqrt{q(k)} \varphi_L(k)$, which proves (21); the $k = 0$ coefficient vanishes. $\square$

Remark 11 (What the L^{-2} collapse does and does not say). The minimum of $q(k)$ over all nonzero finite-volume momenta is $4 \sin^2(\pi/L)$, so a contour isolating the single flat sheet from the other two sheets uniformly over the entire Brillouin zone cannot follow from the second-order symbol. This is an *internal* sheet separation. It does not erase the volume-uniform external electric-shell gap (7), and it does not affect the fixed- k coefficient (20). Neither fact by itself controls the interacting Riesz projection or supplies transfer-matrix overlap; those remain the substantive part of G18.

3.1 The wrapping sheets are cycles, not harmonic — and their harmonic representatives

Theorem 6 exhibits its three H_2 generators explicitly as wrapping sheets $s_{(ij)} = \sum_{x: x_m=0} [x; i, j]$. These satisfy $\partial_2 s = 0$ exactly, so they are cycles and they span the quotient $\ker \partial_2 / \text{im } \partial_3$. They are *not* harmonic: $\partial_3^\dagger s \neq 0$, and

$$\frac{\langle s, \partial_3 \partial_3^\dagger s \rangle}{\|s\|^2} = 2 \quad \text{exactly, for every } L \geq 2. \quad (22)$$

The distinction matters because H_2 is used in two senses: the quotient in Theorem 6, which the sheets span, and the harmonic subspace $\ker \partial_2 \cap \ker \partial_3^\dagger$, which they do not lie in.

Harmonic representatives do exist, and they are one line away. Averaging a sheet over its L translates in the complementary direction,

$$h_{(ij)} = \frac{1}{L} \sum_{r=0}^{L-1} s_{(ij)}(r), \quad \partial_2 h_{(ij)} = 0, \quad \partial_3^\dagger h_{(ij)} = 0, \quad (23)$$

gives three harmonic chains spanning $\ker \partial_2 \cap \ker \partial_3^\dagger$, which is exactly three-dimensional; and $s_{(ij)} - h_{(ij)} \in \text{im } \partial_3$, so a sheet and its average are the same homology class. The averaging is the $k=0$ Fourier projection, and what it removes is precisely the non-harmonic part measured above — which is Remark 9 again, in real space. Nothing about the finding is withdrawn — the generators Theorem 6 exhibits are still not harmonic — but a real-space statement phrased for the harmonic subspace now has representatives to point at, and Proposition 38 covers both, since $B^\dagger w = 0$ holds on all of $\ker \partial_2$.

4 Exact adjacent shared-link dynamics and global assembly

Two plaquettes sharing a link leave six nonshared fundamental links after the external contractions. On the shared link the fusion channels depend on the relative induced orientation,

$$F \otimes \bar{F} = \mathbf{1} \oplus \text{Adj}, \quad F \otimes F = \Lambda^2 F \oplus \text{Sym}^2 F, \quad (24)$$

with dimensions and quadratic Casimirs

$$\begin{aligned} d_{\mathbf{1}} &= 1, & C_{\mathbf{1}} &= 0, \\ d_{\text{Adj}} &= N^2 - 1, & C_{\text{Adj}} &= N, \\ d_{\Lambda^2 F} &= \frac{N(N-1)}{2}, & C_{\Lambda^2 F} &= \frac{(N+1)(N-2)}{N}, \\ d_{\text{Sym}^2 F} &= \frac{N(N+1)}{2}, & C_{\text{Sym}^2 F} &= \frac{(N-1)(N+2)}{N}. \end{aligned} \quad (25)$$

The six nonshared half-links contribute $3C_F$ and the fused link $C_\rho/2$, against an external one-plaquette energy $2C_F$, so the resolvent denominator is $C_F + C_\rho/2$ exactly.

4.1 The channel weights are a theorem

Earlier editions assigned the channel projector squared norm d_ρ/N^2 by an isotropy hypothesis, and every statement in this section rests on that assignment. It is not a hypothesis.

Theorem 12 (Shared-link channel weights). *In the normalised two-plaquette state, the weight of the shared-link channel ρ is exactly d_ρ/N^2 , for both orientation families and every $N \geq 3$.*

Proof. Two steps. First the six nonshared links collapse. Each plaquette contributes a product of three independent Haar links, and a product of independent Haar matrices is Haar, so the amplitude is $\text{Tr}(AU) \text{Tr}(BU^{\pm 1})$ with A, B independent Haar and U the shared link. Integrating A and B contributes δ/N each and leaves a pure degree-(2, 2) moment of U .

Second, two such moments settle both families. In the Weingarten calculus, whose asymptotic origin is [5] and whose finite- N closed forms are [18], the order-two values are $\text{Wg}(e) = 1/(N^2 - 1)$ and $\text{Wg}((12)) = -1/(N(N^2 - 1))$ — the inverse of the S_2 Gram matrix $\begin{pmatrix} N^2 & N \\ N & N^2 \end{pmatrix}$ — the index sums give

$$\begin{aligned} M_{\text{direct}} &= \sum_{ijkl} \langle |U_{ij}|^2 |U_{kl}|^2 \rangle = N^2, \\ M_{\text{cross}} &= \sum_{ijkl} \langle U_{ij} U_{kl} \bar{U}_{il} \bar{U}_{kj} \rangle = N. \end{aligned} \tag{26}$$

The standard formula is often written for $U(N)$. A degree-(2, 2) integrand is invariant under $U \mapsto zU$ for $z \in U(1)$, so integrating it over $SU(N)$ and over $U(N) = (SU(N) \times U(1))/\mathbb{Z}_N$ gives the same number provided the $SU(N)$ -invariants at this degree are exhausted by the permutation contractions — which holds exactly when the degree is below the rank, so that no determinant (Levi-Civita) contraction is available. At degree two this is the condition $N \geq 3$, which is where the construction already lives. For the like family, $\text{Tr}(AU) \text{Tr}(BU)$ decomposes into $\text{Sym}^2 \oplus \Lambda^2$ by symmetrising the two index pairs jointly, and the two squared norms are $(M_{\text{direct}} \pm M_{\text{cross}})/2$. Normalising,

$$n_{\text{Sym}^2} = \frac{N+1}{2N} = \frac{d_{\text{Sym}^2 F}}{N^2}, \quad n_{\Lambda^2} = \frac{N-1}{2N} = \frac{d_{\Lambda^2 F}}{N^2}.$$

For the mixed family the singlet component of $U_{ij} \bar{U}_{lk}$ is $\delta_{il} \delta_{jk}/N$, of squared norm 1 against $M_{\text{direct}} = N^2$, so $n_1 = 1/N^2 = d_1/N^2$ and $n_{\text{Adj}} = 1 - 1/N^2 = d_{\text{Adj}}/N^2$. $\square$

Remark 13. The derivation imports nothing: the Weingarten pair is the inverse of the S_2 Gram matrix and the index sums in (26) are explicit, so the reader can check it without leaving this page. That the pair is not two fitted constants is separately checked — the identity-permutation value times $N^2 - 1$ is 1 identically in N , so the $SU(3)$ numbers are a specialisation of a rank-generic formula. Together these make the *local shared-link coefficient* self-contained. By themselves they do not enumerate every second-order process; that separate step is supplied by Lemma 16. Appendix B records the index sums.

Remark 14 (External corroboration). The same numerator appears in the quantum-simulation literature from an unrelated direction. Ciavarella, Burbano and Bauer [27] publish a finite-rank plaquette matrix element $|M_\rho|^2 = \dim \rho / (\dim A \dim R)$; for two fundamental factors $\dim A = \dim R = N$ and this is d_ρ/N^2 identically in N . That is provenance for a chain already checked here, not a second derivation, and it is not a second derivation. The exhaustion of the shared-link routes is proved independently in Lemma 16.

Remark 15 (A falsified zero backend). The organized authority corpus contains an independent exact diagnostic at $N = 3$:

$$\text{Wg}(e) = \frac{1}{8}, \quad \text{Wg}((12)) = -\frac{1}{24}, \quad \int_{SU(3)} |U_{11}|^4 dU = \frac{1}{6} \neq 0.$$

The related fourth moment $\int |U_{11}|^4 dU = 2/[N(N+1)]$ is $1/6$ at $N = 3$ and in particular nonzero, which is what refuted a claimed Haar vanishing in the balanced (2, 2) sector. It therefore falsifies the archived stranded-flux rule that discards this sector. This is useful negative evidence for retaining the contractions above; it is not a second derivation of t_N .

4.2 The off-diagonal process list is complete

Write $X_q = \chi_q + \bar{\chi}_q$, so $V = \sum_q X_q$, and let $Q_0 = 1 - |0\rangle\langle 0|$.

Lemma 16 (Second-order process completeness). *For $N \geq 3$, $L \geq 3$, and distinct plaquettes p, p' , define*

$$K_{2,-} = P_- V Q_- (E_{F,N} - H_E)^{-1} Q_- V P_-, \quad (27)$$

$$H_{2,-}^{\text{exc}} = K_{2,-} - E_{\text{vac}}^{(2)} P_-, \quad E_{\text{vac}}^{(2)} = \langle 0 | V Q_0 (0 - H_E)^{-1} Q_0 V | 0 \rangle. \quad (28)$$

If p and p' share no link, then $\langle p', - | K_{2,-} | p, - \rangle = 0$. If they share the unique link e , every nonzero intermediate lies in exactly one of $F \otimes \bar{F} = \mathbf{1} \oplus \text{Adj}$, $F \otimes F = \Lambda^2 F \oplus \text{Sym}^2 F$ or the conjugate of the second decomposition. Thus these four shared-link channels exhaust every off-diagonal second-order process.

The conclusion is preserved by the stated link-Casimir budget cutoff, and more generally by cutoffs that are functions only of the link Casimirs and therefore act identically on all four orientation assignments. A general projection inside the degenerate $2E_{F,N}$ eigenspace need not preserve the non-adjacent cancellation.

Proof. Expand the external states and both insertions in oriented characters. A term $\chi_p^\varepsilon \chi_q^\eta$, with $\varepsilon, \eta \in \{\pm 1\}$ and $\chi^{-1} = \bar{\chi}$, has linkwise centre charge $\partial(\varepsilon p + \eta q)$ in $C_1(\Lambda_L; \mathbb{Z}_N)$. Since H_E commutes with the independent link-centre actions, a common electric intermediate between the p and p' sides requires

$$\partial(\varepsilon p + \eta q) = \partial(\varepsilon' p' + \eta' r) \pmod{N} \quad (29)$$

for action plaquettes q, r and signs $\varepsilon, \eta, \varepsilon', \eta'$.

We first record a small-support fact. A nonzero 2-cycle over $\mathbb{Z}_N$ cannot be supported on at most four elementary plaquettes. Choose a supported face f . Each of its four boundary links needs another supported face to cancel its nonzero coefficient, while two distinct cubical faces share at most one link. Four distinct additional faces would therefore be required.

Apply this fact to $z = \varepsilon p + \eta q - \varepsilon' p' - \eta' r$. Equation (29) makes z a 2-cycle over $\mathbb{Z}_N$, so $z = 0$ there. Because $p \neq p'$, the four occurrences must pair with opposite signs. The only possibilities are

$$q = p, \quad \eta = -\varepsilon, \quad r = p', \quad \eta' = -\varepsilon', \quad (30)$$

or

$$q = p', \quad r = p, \quad \eta = \varepsilon', \quad \eta' = \varepsilon. \quad (31)$$

This also closes the determinant exceptions. For $SU(3)$ a coefficient ± 3 would require three occurrences of one face and leave a distinct fourth face with coefficient ± 1 ; for $SU(4)$ a coefficient ± 4 would require all four occurrences to be the same face. Both contradict $p \neq p'$.

The first alternative is absent from the charge-odd vector:

$$X_p |p, -\rangle = \frac{(\chi_p + \bar{\chi}_p)(\chi_p - \bar{\chi}_p)}{\sqrt{2}} |0\rangle = \frac{\chi_p^2 - \bar{\chi}_p^2}{\sqrt{2}} |0\rangle,$$

so its same-face zero- N -ality component, including the vacuum, cancels. Only the exchange alternative remains.

If p, p' are link-disjoint, sharing a vertex being immaterial, linkwise factorisation gives $X_{p'} |p, -\rangle = \sqrt{2} |p, -\rangle |p', +\rangle$ and $X_p |p', -\rangle = \sqrt{2} |p, +\rangle |p', -\rangle$. Both vectors have electric energy $2E_{F,N}$, so the reduced resolvent is scalar on their span, while their inner product is zero by charge orthogonality.

If p, p' share e , the six nonshared links each carry one fundamental or antifundamental factor and the shared link carries exactly two. The two complete tensor-product decompositions above therefore give all common intermediates. Their energies and gaps are $E_\rho = 3C_F + \frac{1}{2}C_\rho$ and $E_\rho - E_{F,N} = C_F + \frac{1}{2}C_\rho > 0$. Spectral resolution yields only their four projectors. Vacuum subtraction is diagonal and leaves this cross matrix element unchanged. The two-face term is connected, so the linked Möbius projection leaves it unchanged as well. $\square$

With Theorem 12 the resolvent weight is

$$w_\rho = -\frac{d_\rho/N^2}{C_F + C_\rho/2}, \quad (32)$$

and summing within the two orientation families,

$$A_N = w_{\mathbf{1}} + w_{\text{Adj}} = -\frac{2N^3}{(N^2 - 1)(2N^2 - 1)}, \quad (33)$$

$$B_N = w_{\Lambda^2} + w_{\text{Sym}^2} = -\frac{4N(N^2 - 2)}{(N^2 - 1)(4N^2 - 9)}. \quad (34)$$

4.3 The shared-link law

Theorem 17 (All-rank adjacent shared-link law). *Let $p \neq p'$ be oriented plaquettes sharing the single link e . For every integer $N \geq 3$, let $s = s(p, e)$, $s' = s(p', e)$, and let $\hat{\Pi}_\rho$ be the charge-conjugation-closed electric projector onto the shared-link channel ρ (including the conjugate irrep on the conjugate like-channel subspace). Define*

$$\eta_\rho = \begin{cases} -1, & \rho \in \{\mathbf{1}, \text{Adj}\}, \\ +1, & \rho \in \{\Lambda^2 F, \text{Sym}^2 F\}. \end{cases}$$

Then

$$\hat{\Pi}_\rho X_{p'} |p, -\rangle = \eta_\rho s s' \hat{\Pi}_\rho X_p |p', -\rangle, \quad \|\hat{\Pi}_\rho X_{p'} |p, -\rangle\|^2 = \frac{d_\rho}{N^2}, \quad (35)$$

$$\langle p', - | X_p \hat{\Pi}_\rho X_{p'} |p, -\rangle = \eta_\rho s(p', e) s(p, e) \frac{d_\rho}{N^2}. \quad (36)$$

Consequently the linked order- u^2 contribution through all four channels is

$$\langle p', - | H_2 | p, -\rangle = t_N s(p', e) s(p, e) = t_N (BB^\dagger)_{p'p},$$

where $s(p, e) = \pm 1$ is the cellular boundary incidence and

$$t_N = B_N - A_N = \frac{2N(N^2 - 4)}{(N^2 - 1)(2N^2 - 1)(4N^2 - 9)} > 0.$$

Proof. Because distinct cubical faces share at most one link, $(BB^\dagger)_{p'p} = ss'$. Expand $|p; \varepsilon\rangle_{\text{or}} := \chi_p^\varepsilon |0\rangle$, with $\chi_p^{+1} := \chi_p$ and $\chi_p^{-1} := \bar{\chi}_p$, so that $|p, -\rangle = 2^{-1/2} \sum_{\varepsilon=\pm 1} \varepsilon |p; \varepsilon\rangle_{\text{or}}$. On the shared link the two oriented factors are εs and $\varepsilon' s'$. A like channel requires $\varepsilon \varepsilon' = ss'$, while a mixed channel requires $\varepsilon \varepsilon' = -ss'$. Thus the two projected route vectors differ by exactly $\eta_\rho ss'$. Their two charge-conjugate assignments are orthogonal for $N \geq 3$; the two factors $1/\sqrt{2}$ therefore sum to one channel weight, not twice one. Theorem 12 gives the norm d_ρ/N^2 , proving (35) and (36).

On this channel, $(E_{F,N} - H_E)^{-1} \hat{\Pi}_\rho = -\hat{\Pi}_\rho / (C_F + C_\rho/2)$. Its contribution is therefore $\eta_\rho ss' w_\rho$. Summing gives $ss' (w_{\Lambda^2 F} + w_{\text{Sym}^2 F} - w_{\mathbf{1}} - w_{\text{Adj}}) = ss' (B_N - A_N)$. Substitution of (33)–(34) gives the displayed rational function. Every denominator factor and $N^2 - 4$ is positive for $N \geq 3$. $\square$

The opening geometric ingredient of the proof is no longer prose: the verifier enumerates every distinct face pair of the periodic complex at $L = 3$ and $L = 4$ and confirms that no pair meets in more than one link. Two faces sharing a link span at most two sites per axis, so these windows already exhibit every local configuration and larger L adds translates only.

Corollary 18 (Global torus assembly). *On the trivial-flux charge-odd electric shell of Theorem 1, translation and cubic symmetry make the remaining diagonal contribution a scalar, and*

$$H_{\text{eff},-}^{(N)}(k, u) = E_{\text{flat},N}(u)I + t_N u^2 B(k)B(k)^\dagger + O(u^3).$$

Consequently the three bands through second order are $E_{\text{flat},N}(u)$ and $E_{\text{flat},N}(u) + t_N u^2 q(k)$ twice, and the flat carrier is $\ker \partial_2$, of dimension $L^3 + 2$.

Proof. The centre-charge argument of Lemma 16, now with only three plaquette occurrences, gives $\langle p', -|P_- V P_-|p, - \rangle = 0$ for $p \neq p'$. A possible $SU(3)$ determinant balance would require all three occurrences to be the same face and is therefore excluded off diagonal. Translation and cubic symmetry make the remaining first-order diagonal a scalar, including any $SU(3)$ same-face determinant contribution, and it is absorbed into $E_{\text{flat},N}(u)$.

At second order Lemma 16 and Theorem 17 assemble every off-diagonal entry. Translations act transitively on faces of a fixed orientation and the cubic group permutes the orientations, so the invariant second-order diagonal is also a multiple of the identity. Thus the adjacency part is $t_N(BB^\dagger - 4I)$; absorb $-4t_N$ into the scalar and apply Theorem 5. $\square$

The status change deserves one sober sentence rather than celebration. What earlier editions carried as two prose premises is now proved relative to the standard degenerate perturbation framework of (18), the classical Casimir arithmetic of Theorem 1, and the centre-charge classification of Lemma 16. The finite-volume $L = 3, 4, 5$ enumerations in the verifier are regressions on those proofs, not substitutes for them. The charge-even identification of Section 6 remains a hypothesis, and nothing about convergence, higher orders, or the continuum has changed.

Two consequences,

$$t_3 = \frac{5}{612}, \quad t_N = \frac{1}{4N^3} - \frac{1}{16N^5} - \frac{77}{64N^7} - \frac{1021}{256N^9} + \dots \quad (37)$$

The expansion is governed by a closed form. Writing the planar-normalised coefficient as $4N^3 t_N$,

$$1 - 4N^3 t_N = \frac{2N^4 + 31N^2 - 9}{(N^2 - 1)(2N^2 - 1)(4N^2 - 9)}, \quad (38)$$

an identity of rational functions, and the right side is positive and the left strictly decreasing for every real $N \geq 2$: substituting $N = M + 2$ makes every numerator coefficient of both the remainder and the derivative nonnegative, so $4N^3 t_N$ increases monotonically to 1. At fixed canonical coordinate u the conditional bandwidth of Corollary 18 is $12t_N u^2$, so the entire dispersive content of the projected sector switches off at exactly the rate N^{-3} , with computable subleading coefficient $-1/(4N^2)$: conditional on the assembly, the carrier's exact degeneracy is approached by the *whole* sector in the planar limit. This is consistent with the general suppression of glueball mixing at large N , but no 't Hooft-coupling identification is made here: the statement is about the coefficient in the canonical coordinate, and it is exact.

4.4 The vacuum-mediated route

Lemma 16 is specific to the charge-odd shell. The corresponding assertion is false in general: the charge-even sector admits a vacuum-mediated route. Its cancellation in the odd sector is the charge-conjugation selection rule below, and omitting it is a recorded error, not a hypothetical one.

Proposition 19 (Vacuum-mediated route). *The reduced-resolvent route*

$$\frac{\langle p'|V|0\rangle\langle 0|V|p\rangle}{E_{F,N} - E_0} = \frac{\langle p'|V|0\rangle\langle 0|V|p\rangle}{E_{F,N}}$$

is available for any pair of plaquettes, sharing a link or not. It vanishes identically in the charge-odd sector, because V is charge-even and $|0\rangle$ is charge-even, so $\langle 0|V|p, -\rangle = 0$. Consequently the charge-odd hopping $t_N = B_N - A_N$ carries no such term, while the charge-even shared-link element acquires $1/C_F$: the normalised charge-even matrix-element product is 2 and $E_{F,N} = 2C_F$.

$$\ell_N = A_N + B_N + \frac{1}{C_F} = -\frac{2N(3N^2 - 5)}{(N^2 - 1)(2N^2 - 1)(4N^2 - 9)}.$$

Remark 20. At $N = 3$, $A_3 + B_3 = -481/612$ and $1/C_F = 3/4$, giving $\ell_3 = -11/306$. The first of these is exactly the value once published for the charge-even hopping and later superseded; the correction is precisely the omitted route. Stating the selection rule turns a one-rank erratum into an all-rank formula.

In the charge-odd shell, Lemma 16 and this proposition dispose of non-adjacent pairs: the exchange vectors are orthogonal and the vacuum route vanishes. In the charge-even sector, however, this proposition does not dispose of the non-adjacent vacuum route. Taken over all pairs it would be a rank-one operator supported at $k = 0$, not the twelve-neighbour adjacency multiplied by ℓ_N . The check above establishes the all-rank shared-link value ℓ_N ; the reorganisation or cancellation of the charge-even all-pairs term remains a separate requirement.

4.5 Finite-volume width

Let $q_{\max}(L)$ be the largest sampled q on the periodic L^3 grid. Then

$$q_{\max}(L) = \begin{cases} 12, & L \text{ even,} \\ 12 \cos^2(\pi/2L), & L \text{ odd,} \end{cases} \quad (39)$$

because only at even L does some k_j reach π exactly. The width of the full retained charge-odd manifold is $W_{N,L}^{(-)} = q_{\max}(L) t_N u^2 + O(u^3)$, so for even L , and coefficientwise as $L \rightarrow \infty$ at fixed perturbative order, $[u^2]W^{(-)} = 12t_N \sim 3/N^3$. This concerns the stiffness of the dispersive doublet, not motion of the lowest branch, which is flat at this order.

The smallest nonzero incidence eigenvalue is $q_{\min} = 4 \sin^2(\pi/L)$, attained by one unit of momentum in a single direction.

4.6 The orientation signs are essential

Removing the orientation signs — from the entries of (19) and from the phases, so each $e^{ik} - 1$ becomes $e^{ik} + 1$ — gives an unsigned incidence $\mathcal{N}(k)$. Define its comparison adjacency by $A_+(k) := \mathcal{N}(k)\mathcal{N}(k)^\dagger - 4I$. Its adjacency spectra at Γ and R are $\{12, 0, 0\}$ and $\{-4, -4, -4\}$, which share no




Figure 2: Order- u^2 coefficients along Γ - X - M - R - Γ . The red carrier and dispersive doublet are the signed incidence result; their conditional global assembly is Corollary 18. The blue curves are the exact unsigned-incidence comparison spectrum. Reading the blue curves as physical charge-even bands requires the separate symbol hypothesis of Section 6.

eigenvalue, so no level can be momentum independent. The flat branch is therefore not a generic consequence of three face orbitals or cubic symmetry.

Section 6 computes the unsigned comparison spectrum in closed form. The signed span is $8 - (-4) = 12$, giving the charge-odd second-order width $12t_3 = 5/51$ at $N = 3$. The unsigned continuous-zone span is $12 - (-4) = 16$. Interpreting $16|\ell_3| = 88/153$ as a charge-even bandwidth is conditional on deriving the complete linked charge-even symbol; on a finite torus the value -4 is sampled only when L is even.

5 The second-order coefficient on a two-cube space

Everything above derives t_N from a resolvent ledger indexed by the fusion channel of one shared link. That ledger is representation theory carried out on a single plaquette pair; it has never been tested against an operator calculation on a space large enough for two cubes to interact. This section reports such a test and, more importantly, states what it settles.

The construction is the open face-sharing $(3, 2, 2)$ prism with its eleven oriented plaquettes, in the exact-Casimir $\mathcal{B} = 6$ $SU(3)$ truncation, of basis dimension 1,590,462. The delivery reports that it reaches the relevant sector through 794 paths and 398 reachable states; it does not assemble a 1,590,462-dimensional magnetic matrix. Its complete charge-odd one-plaquette shell at $E_* = 8/3$ has dimension eleven — a census result and not the plaquette count. The full E_* eigenspace has dimension 22, split $11_{C=+} \oplus 11_{C=-}$ by tensor-derived charge conjugation, and the delivery removes all 22 from the pseudoinverse; because the eleven columns exhaust the odd half, that agrees with the $Q = 1_- - P_-$ prescription of (18). Writing K_{LR} for the raw second-order operator and K_L, K_R, K_F for the left-cube, right-cube and shared-face operators in their own coordinates, the connected

operator is defined by the literal Möbius fold

$$K_{\text{conn}} = K_{LR} - J_L K_L J_L^\dagger - J_R K_R J_R^\dagger + J_F K_F J_F^\dagger, \quad (40)$$

and the same fold applied to the geometry gives G_{conn} . The fold is literal in the sense that no eigenvalue or fitted scalar is subtracted: each source operator is folded in its own coordinates and lifted. The delivery corroborates that internally by applying the same weights history by history over its 1,984 ordered two-step histories, of which 1,192 cancel at weight zero, and the two routes agree to 4.9×10^{-15} .

First order settles itself. On the shell $PVP = I_{11}$, so the first-order term is nonzero but scalar: it shifts all eleven branches equally and creates no branch ambiguity. Its connected fold then vanishes *identically*, and by counting rather than by cancellation — each of the eleven faces is covered exactly once by $\{L, R\}$ once the shared face is restored by F , so every Möbius weight in (40) applied to the identity is $1 - 1 = 0$. Second order is therefore the leading connected term and not a correction to a first-order one.

At second order the result is

$$K_{\text{conn}}^{(2)} = \frac{5}{612} G_{\text{conn}} + D, \quad (41)$$

with D diagonal and delivered. The geometry alone fixes the *splitting* inside each block and nothing else: G_{conn} turns out to be four disjoint 2×2 blocks plus three isolated directions (Appendix F), so each pair splits by $2 \times 5/612 = 5/306$ about its own diagonal entry. Given D , which Appendix F prints, the spectrum is then closed form: 0 once, $-\frac{15}{4}$ twice, $-\frac{34}{9}$ four times, and $-\frac{129}{34}$ four times — exact arithmetic on entries that are themselves reconstructions, a boundary made precise below. The delivery reports the coefficient recovered target-blind — supplied to no builder, graph fit, branching rule or validator — and reports that the payload hashes were produced before the recorded comparison, a chronology it itself calls record-backed and only partly auditable; the value is byte-present in the nested $\mathcal{B} = 4$ package that the $\mathcal{B} = 6$ build seals as an authority object, so the defensible statement is dependency-path nonuse and not corpus-wide absence. This paper re-thresholds those delivered gates rather than re-running the build, including a deliberately wrong-sign control that the held-out scaling rejects: on a 66-dimensional held-out word space the correct second-order residual falls with log-log slope 2.90 while the sign-reversed control falls with slope 2.00, a minimum error ratio of 325. What all this establishes is runtime nonuse of the comparator on the delivered dependency path, which is not a preregistration. The finite- u validations are asymmetric: the $\mathcal{B} = 4$ delivery includes a held-out diagonalisation of the full 4,180-dimensional charge-odd space with an $O(u^3)$ residual, whereas the $\mathcal{B} = 6$ check is restricted to the 66-dimensional $P + Q_1$ star.

Resolving (41) by the irreducible representation the shared link actually carries in the ranked intermediate state gives six coefficients rather than four, because a link distinguishes ρ from $\bar{\rho}$:

$$\begin{array}{c|cccccc} \rho & \mathbf{1} & \mathbf{3} & \bar{\mathbf{3}} & \mathbf{6} & \bar{\mathbf{6}} & \mathbf{8} \\ \hline c_\rho & \frac{1}{12} & -\frac{1}{12} & -\frac{1}{12} & -\frac{1}{9} & -\frac{1}{9} & \frac{16}{51} \end{array} \quad (42)$$

That these six sum to $5/612$ is weak evidence: six rationals reach one target in many ways, and a fitted set would pass it. The statement worth making is stronger, and it is what makes this more than a coincidence of one sum. First, though, the boundary.

The six coefficients of (42) are not symbolic: the delivery stores them as rational reconstructions of pinned finite-precision Clebsch–Gordan contractions, with a maximum reconstruction residual of 2.2×10^{-15} per channel — 2.1×10^{-14} across the delivery’s reconstructed matrices as a whole — and its own certificate declaring `symbolic_exact_local_amplitudes: false`. What follows

is therefore reported, conditional on those six numbers, and not a symbolic identity in the sense of Theorem 12.

Reported computational result 21 (The census is the Weingarten ledger). At $N = 3$ the six reconstructed link-irrep coefficients of (42) are the four fusion-channel weights (32) individually:

$$\begin{aligned} c_1 &= -w_1, & c_8 &= -w_{\text{Adj}}, \\ c_3 &= c_{\bar{3}} = \tfrac{1}{2}w_{\Lambda^2 F}, & c_6 &= c_{\bar{6}} = \tfrac{1}{2}w_{\text{Sym}^2 F}. \end{aligned}$$

Substituting (25) into (32) at $N = 3$ gives $w_1 = -1/12$, $w_{\text{Adj}} = -16/51$, $w_{\Lambda^2 F} = -1/6$ and $w_{\text{Sym}^2 F} = -2/9$; comparison with (42) is then arithmetic. Six reconstructed rationals landing on four predicted ones, with the two coincidences the correspondence requires, is a strong consistency statement and not a proof that the underlying contraction is exact.

One further feature of the census is worth separating out, because it is the paper’s central architecture measured rather than assumed. The six link-resolved matrices are not merely six numbers that sum correctly: each is *separately* proportional to the same incidence Gram. The eleven oriented plaquettes of the prism have exactly 56 ordered adjacent pairs, each sharing one link with signed overlap ± 1 ; the delivery reports a constant channel-to-geometry ratio on all 56 of them, for every one of the six channels, with worst residual 2.2×10^{-15} . Colour dynamics fixes a number per channel and the cellular boundary operator fixes the matrix; if the two mixed, the ratio would be constant for the total and not for the parts.

The correspondence is scoped to what the Casimir closure exhausts: adjacent, one-action, shared-link intermediates at order u^2 . Same-face diagonal routes, on-site terms and higher-action intermediates lie outside it — the $\mathcal{B} = 6$ scalar below is exactly a place where one of them is missing — and so does any untruncated-Hilbert-space statement.

It is scoped in one further way, which matters more. The two-cube build draws its local Wilson amplitudes from the same hash-pinned upstream Clebsch–Gordan archive as the single-cube reconstructions of the next subsection, and the squared amplitudes that archive produces are the published d_ρ/N^2 [27] — the numerator of (32). The census therefore does not confirm the weights from an independent source: it tests the *assembly* that carries them, which is the reachable-channel list, the resolvent denominators, the operator-level fold, the even split between conjugates, and the geometry that survives each channel separately. That is the statement worth making, and it is a different statement from confirmation.

Two features of the correspondence carry the content. The mixed family enters negated because $t_N = B_N - A_N$, which is the same relative sign the incidence orientation carries. And each like-family *fusion* weight is split evenly between the irrep and its conjugate — the two orientations of the shared link weighted equally — which the delivery resolves on the two-cube space rather than assuming. It is not measured *here*: what this repository checks is the identity the six delivered numbers satisfy.

Conjecture 22 (All-rank even split). For every $N \geq 3$ each like-family fusion weight splits evenly between an irrep and its conjugate in the link-resolved census.

This is a conjecture and nothing here promotes it. Falsifying it means a link-resolved census at some $N \geq 4$ in a truncation that still retains every adjacent shared-link channel; at $N = 4$ the endpoint budgets make that $\mathcal{B} \geq 33/4$, and no such construction has been run.

5.1 One truncation, two constructions

The two-cube space also isolates why a cutoff can reverse the sign of t_3 . An adjacent shared-link channel ρ is reachable only if the endpoint Casimir budget $C_2(\rho) + 2C_2(\mathbf{3})$ fits inside the truncation, meaning $C_2(\rho) + 2C_2(\mathbf{3}) \leq \mathcal{B}$: the sextet needs exactly 6 and the adjoint 17/3, so both survive $\mathcal{B} = 6$ and neither survives $\mathcal{B} = 4$.

The weak inequality is not a convenience. Both retentions this section uses are decided by an equality: $\Lambda^2 F = \mathbf{3}$ has budget exactly 4 and so sits exactly on the $\mathcal{B} = 4$ cutoff, and the sextet sits exactly on the $\mathcal{B} = 6$ one. Under a strict rule $\mathcal{B} = 4$ would reach $\{\mathbf{1}\}$ alone and $\mathcal{B} = 6$ would lose the sextet, contradicting both delivered censuses — so the deliveries are what fix the convention, which is also why the budget arithmetic is not an independent confirmation of either census.

Corollary 23 (What a link cutoff drops). *Retaining only the singlet and $\Lambda^2 F$ routes projects the second-order hopping to*

$$w_{\Lambda^2 F} - w_{\mathbf{1}} = -\frac{1}{12},$$

opposite in sign to t_3 and 10.2 times larger in magnitude. What is dropped is

$$w_{\text{Sym}^2 F} - w_{\text{Adj}} = \frac{14}{153},$$

and $-\frac{1}{12} + \frac{14}{153} = \frac{5}{612} = t_3$. This is arithmetic on (32) alone.

The published $p + q \leq 1$ link cutoff retains exactly those two routes — a per-link restriction, where $\mathcal{B} = 4$ is a total endpoint budget, so the two schemes agree on the retained channels at this order and are not the same rule — so $-1/12$ is its projected adjacent hopping and $14/153$ is the adjacent-hopping correction that completes it — not $5/612$, which would double-count what it already carries. On the two-cube space the first of those two numbers is delivered rather than inferred: a separately sealed $\mathcal{B} = 4$ run on the identical geometry gives $K_{\text{conn}}^{(2), \mathcal{B}=4} = -\frac{1}{12}G_{\text{conn}} + D_{\mathcal{B}=4}$ outright, and the Casimir budgets of the previous paragraph say without reference to any census that $\mathcal{B} = 4$ can reach only $\{\mathbf{1}, \mathbf{3}, \bar{\mathbf{3}}\}$. Conditional on (42), the same split of the census reproduces it, $c_1 + c_3 + c_{\bar{3}} = -1/12$, and gives the omitted half as $c_6 + c_{\bar{6}} + c_8 = 14/153$.

Remark 24. Two single-cube reconstructions bracket the same statement. A full $\mathcal{B} = 4$ open cube assembled from the pinned `ymcirc` $d = 3$, $\mathcal{B} = 4$ plaquette data [28, 29] reproduces $-1/12$ and the truncated shell ordering to 4.3×10^{-11} ; a $\mathcal{B} = 6$ cube reproduces $+5/612$ and the inverse ordering to 8.4×10^{-15} . Both locate the missing weight at the same $14/153$.

Neither is an independent replication. Both read tabulated plaquette matrix elements of the same lineage — and the two-cube build draws its Clebsch–Gordan amplitudes from the same hash-pinned `pyclebsch` archive as the $\mathcal{B} = 6$ cube, so that agreement is a cross-check of assembly and not of the amplitudes. The published closed form [27] is d_ρ/N^2 — the very numerator of (32). This is a cross-check of assembly, not a third confirmation of t_3 ; the exact rational identity is Corollary 23, which needs no tolerance. The $p + q \leq 1$ nomenclature and the finite-rank matrix element are both from [27], and single-cube $SU(N)$ diagonalisation in this style goes back to [17].

The correction to the adjacent-hopping coefficient of a published T_1 Hamiltonian is therefore $14/153$ and not $5/612$, which would double-count the two routes T_1 already carries. One coincidence of value is worth naming: $-1/12$ is also the singlet weight $w_{\mathbf{1}}$, and the cutoff value is a *difference* of two weights, not a weight.

The $\mathcal{B} = 6$ cube is channel-complete for the hopping and *not* for the scalar: its scalar is $39/68$ against the bridge’s $11/34$, differing by exactly $1/4$, the local same-face sextet route, which first enters at cutoff $20/3$ — integer label $\mathcal{B} = 7$, so any cutoff above $20/3$ already carries it.

That is a prediction with a cheap test, and the test was executed. A six-face $\mathcal{B} = 7$ reachable-state probe returns $\alpha_{\mathcal{B}=7} = 11/34$ and $t_{\mathcal{B}=7} = 5/612$ with scalar-plus-Gram residual 1.1×10^{-15} , so raising the cutoff past the same-face threshold lowers all three absolute coefficients by exactly $1/4$ and leaves the relative shell $\{0, 5/153, 5/102\}$ untouched — the scalar moves, the hopping does not. That probe’s artifact did not travel, so it is reported here and not checked.

5.2 What the cutoff moves, and where

The connected diagonal D is eleven numbers in both deliveries and is treated as a remainder in both. It is not eleven numbers. The open prism has a symmetry group — the cell exchange $x \mapsto 2 - x$, the two transverse reflections and the $y \leftrightarrow z$ swap, of order 16 — and D is constant on its orbits, which are $8 + 2 + 1$: the eight side faces, the two end caps, the shared face.

Proposition 25 (Where the truncation lives). *The eight-element orbit is exactly the support of the four cross-cell pairs, and it is the only orbit whose diagonal value differs between the two truncations: the end caps carry $-15/4$ and the shared face 0 at both $\mathcal{B} = 4$ and $\mathcal{B} = 6$, while the side value moves from $-7/4$ to $-2317/612$. Together with the four cross-cell blocks carrying the whole off-diagonal change, the entire $\mathcal{B} = 4 \rightarrow \mathcal{B} = 6$ difference is confined to the faces that carry connected transport.*

This also says what the nonuniform D is. Where the faces of a complex form a *single* orbit — the closed cube of the next subsection, or the periodic torus of Corollary 18 — orbit-constancy forces the diagonal to be scalar, and the operator is exactly $\alpha I + tG$. The prism has three orbits because it has a boundary. D is the open boundary showing up in the diagonal, not a failure of the incidence form.

5.3 The one-cube shell, and why its flat level cannot move

Those two single-cube reconstructions are usually read as numerical cross-checks. They are more than that, because the object they diagonalise is the same chain complex. The six faces of a cube are a closed surface, so in the outward orientation $\ker \partial_2$ is one-dimensional and spanned by the sum of all six — the fundamental class, $b_2(S^2) = 1$. That is Theorem 6 on a complex with no three-cells, where the $L^3 - 1$ boundaries are absent and only the class survives.

Proposition 26 (One-cube shell). *On the closed cube surface $G = \partial_2^T \partial_2$ commutes with all 24 proper rotations. The combinations $b_{+i} - b_{-i}$ carry the defining vector representation and the combinations $b_{+i} + b_{-i}$ the permutation representation of the three unordered axes, so with the charge-odd inversion $Pb_{+i} = -b_{-i}$ the shell is*

$$A_1^{--} \oplus T_1^{+-} \oplus E^{--},$$

with G -eigenvalues 0, 4, 6 and multiplicities 1, 3, 2. Consequently $\alpha I + tG$ has levels α , $\alpha + 4t$, $\alpha + 6t$.

Proof. The rotations permute the six outward faces, and G is built from the boundary map, so it commutes with the permutation action. On the antisymmetric combinations that action is the defining representation by inspection — $b_{+i} - b_{-i}$ transforms as e_i — and on the symmetric combinations it is the permutation representation of $\{\pm e_i\}$, which is $A_1 \oplus E$. Under $Pb_{+i} = -b_{-i}$ the antisymmetric sector is parity even and the symmetric one parity odd. In these sectors G is $4I$ and $4I - 2(J - I)$ respectively, of spectra $\{4, 4, 4\}$ and $\{0, 6, 6\}$; since the three isotypic components have distinct dimensions the eigenvalues attach to them uniquely. $\square$

Three things follow. The cubic labels are derived here rather than assigned: both deliveries assert them and neither builds the symmetry matrices. The A_1^- level is Proposition 38’s mechanism on a second complex — $Gw = 0$ there, so that level sits at α for every t , and no truncation, no channel and no coefficient in dispute anywhere in this paper can move it. And the reversal the two runs report is therefore a statement about the sign of one number rather than about the shell as a whole: at $\mathcal{B} = 6$ the levels are $\{39/68, 371/612, 127/204\}$ and at $\mathcal{B} = 4$ $\{37/12, 11/4, 31/12\}$, in the opposite order, with A_1^- in each case sitting exactly at the scalar. The relative shell $\{0, 4t, 6t\}$ is $\{0, 5/153, 5/102\}$ at $\mathcal{B} = 6$.

The finite-volume fingerprint of the 28 August finite- N bridge certificate — a different release, and named here because a reader tracing it to the two-cube one would find nothing — is worth recording because its most discriminating multiplicity is not a feature of the truncation: on the periodic 3^3 torus the incidence eigenvalues 0, 3, 6, 9 have multiplicities 29, 12, 24, 16, respectively, and $29 = L^3 + 2$ is Theorem 6.

Remark 27 (Scope of the second-order delivery). Everything above in this section is second order in u , finite volume, strong coupling and charge odd. It bears on t_N and decides nothing about the fourth-order coefficient of Section 9. The $\mathcal{B} = 6$ two-cube construction was not re-run here: its arithmetic is audited against geometry rebuilt independently, and its diagonal D is B -truncated and not proved cutoff-stable. The older held-out finite- u validation that rejects the wrong-sign control is a 66-dimensional $P + Q_1$ word space with the Q_1 - Q_1 magnetic block set to zero; it is not the radius-two $K_{2,\star}$ operator defined below. The two branches are not equally covered: at $\mathcal{B} = 4$ the delivery reports a direct diagonalisation of the complete 4,180-dimensional charge-odd block at four small couplings, two of them held out, with the residual after $E_* + u + u^2 \text{ eig } K_{\text{eff}}^{(2)}$ scaling as u^3 ; at $\mathcal{B} = 6$ no such diagonalisation exists and the 66-dimensional star stands in for it. That $\mathcal{B} = 4$ report carries no check here — its certificate did not travel — and is recorded as the delivery’s rather than as this paper’s. No external group has reproduced that release. The later radius-two release is a distinct construction. It neither upgrades the $\mathcal{B} = 6$ cutoff-stability claim nor supplies a radius-three test.

5.4 A detached-replayed radius-two Ritz extension

A later sealed release (29 August 2026) on the same open $(3, 2, 2)$ prism assembles the charge-odd and charge-even $SU(3)$ two-cube Ritz operators with no Casimir cutoff inside its declared frontiers: the electric seed P (eleven charge-odd faces), the complete one-action frontier Q_1 , and the exact-energy-factorized two-action frontier Q_2 , with cumulative charge-odd dimensions 11, 77, 773 and charge-even dimensions 1, 2, 9. The release calls this construction *cutoff-free*: here that means no $\mathcal{B}$ cutoff inside the declared one- and two-action frontiers. The radius-two projection itself remains a truncation of the gauge-theory Hilbert space, and the phrase is used in no wider sense below.

To keep action radius distinct from perturbative order, write the three nested matrices as K_1 , $K_{2,\star}$ and $K_{2,\text{full}}$. In $P + Q_1 + Q_2$ block order,

$$H_0 = \text{diag}(E_s I, H_1, H_2), \quad V = \begin{pmatrix} A_0 & B & 0 \\ B^\dagger & G & C \\ 0 & C^\dagger & D_2 \end{pmatrix}. \quad (43)$$

Here $E_s = 8/3$ in the odd seed and $E_s = 0$ in the even seed. The three finite-coupling operators

use $K(u) = H_0 + uV$, with

$$\begin{aligned} K_1(u) &= (H_0 + uV)|_{P+Q_1}, \\ K_{2,\star}(u) &= H_0 + u[V \text{ with only } D_2 = Q_2 V Q_2 \text{ set to zero}], \\ K_{2,\text{full}}(u) &= (H_0 + uV)|_{P+Q_1+Q_2}. \end{aligned} \tag{44}$$

Thus “full” means complete only within the retained Q_2 block. Every finite-coupling Cauchy interlacing inequality between the three is checked, with recorded violation exactly 0. All quoted quantities are invariant under diagonal basis rephasing, which is why no signed hopping entry is read as an observable there. The delivery is finite-precision throughout; everything in this subsection is reported at stated tolerances, never promoted.

Reported computational result 28 (The rational sub-block of the cutoff-free second-order spectrum). Six of the eleven eigenvalues of the sealed cutoff-free second-order block $C_2 = BR_1 B^\dagger$ on the odd seed are rationals to within 10^{-12} (the worst residual is 5×10^{-14} , about 30 ulps):

$$-\frac{2429}{306} (\times 2), \quad -\frac{404}{51} (\times 3), \quad -\frac{2419}{306} (\times 1),$$

and the outer levels straddle the central triple by exactly

$$\frac{2429}{306} - \frac{404}{51} = \frac{404}{51} - \frac{2419}{306} = \frac{5}{306} = 2t_3.$$

The remaining five eigenvalues are not rational at this precision.

The point is what this operator was *not* told. The census of Reported result 21 recovered the shared-link channel coefficients from a build that retained those channels; here the sealed radius-two operator retains every second-order route the two-action frontier supports, makes no shared-link truncation, and the shared-link coefficient $t_3 = 5/612$ of Theorem 17 appears anyway — as the exact splitting scale of a rational sub-block, embedded in a spectrum whose other levels are algebraic. A structural identification of that sub-block (which face combinations carry it, and why the central level is threefold) is not derived here and is recorded as a lead, not a claim; the open two-cube cluster has face-inequivalent environments, so a non-scalar diagonal is expected and no contradiction with Corollary 18, which is a torus statement, arises.

The nested-difference coefficients have a provable onset, which is what makes them readable at all.

Proposition 29 (Onset of the nested-radius differences). *Let $R_i = (E_s I - H_i)^{-1}$. If $PVQ_2 = 0$, the leading seed-space operator present in $K_{2,\star} - K_1$ is*

$$I_4 = BR_1 CR_2 C^\dagger R_1 B^\dagger, \quad I_5 = BR_1 CR_2 D_2 R_2 C^\dagger R_1 B^\dagger, \tag{45}$$

where I_5 is the leading term present in $K_{2,\text{full}} - K_{2,\star}$.

Proof. Any contribution that enters Q_2 must follow the shortest closed block walk $P \rightarrow Q_1 \rightarrow Q_2 \rightarrow Q_1 \rightarrow P$, which contains four insertions of V . Distinguishing the full operator from the star operator additionally requires one insertion of D_2 , hence five insertions. The blocks A_0 and G may dress these walks but cannot shorten them. Evaluating the intermediate resolvents at the seed energy gives (45). $\square$

The proposition is exact conditional on the declared frontier and sealed block structure. The following values are direct floating-point evaluations of its closed products, not exact rationals and not eigenvalue extrapolations.

Reported computational result 30 (Detached radius-two replay). The lowest eigenvalue of the odd second-order seed operator is isolated by 0.08438094122012973; the even seed is one-dimensional. Projection of (45) gives

	$C = -$		$C = +$	odd-even gap
I_4	-25.63495194947556	-24.807728943850304	-0.8272230056252567	
I_5	-4.933069330883140	-9.116955631684514	+4.183886300801373	

The last entry is the leading $K_{2,\text{full}} - K_{2,\star}$ contribution to the retained odd-even gap. It is *not* the complete fifth-order coefficient of either unablated sector Hamiltonian. The weak- u cubic intercept 4.183884987168061 differs by 1.31×10^{-6} ; the direct product, not that fit, defines the reported coefficient.

The hash-pinned detached verifier starts from both sealed parity-sector parent operators and recomputes the two coupling grids, all three nested spectra, residuals, interlacing, direct products, rephasing and adaptive path integrals. It returned zero maximum recomputation error and maximum rephasing discrepancy 4.09×10^{-14} ; all 24 hostile, algebra, numerical-safety and serialization tests passed. This is an independent downstream replay of the sealed parents, not an independent derivation of their microscopic gauge-basis contractions.

The boundary is quantitative. No odd Q_3 matrix is present, so no matched radius-two error enclosure or convergence claim exists. At $g = 1$ the odd and even ground seed weights have fallen to 0.274 and 0.277, while the even exterior residual is 3.784; the retained eigenvalues are still accurately computed, but their accuracy as approximations to a larger space is uncontrolled there.

One further feature bears on Section 9. The delivery's fifth-order object on the eleven-dimensional seed is recorded as a *matrix*: its traceless part is nonzero, with $\text{tr}(\text{shape}^2) = 0.359$ and eleven recorded shape eigenvalues. It is therefore the first delivered instrument in this corpus whose output is not a Γ -point scalar. That observation is about the instrument class only: an open prism carries no off-axis momentum, so I_4 and I_5 in (45) do not choose a value of C_{shp} , and nothing here weakens Theorem 40.

6 The unsigned-incidence comparison operator

The unsigned incidence $\mathcal{N}(k)$ of Section 4 defines an exact algebraic comparison operator $G_+(k) = \mathcal{N}(k)\mathcal{N}(k)^\dagger$. Its identification with the complete linked charge-even effective Hamiltonian is an additional hypothesis, not a result of the shared-link calculation: Proposition 19 exposes an all-pairs vacuum route whose cancellation or linked-cluster reorganisation has not been derived. With that boundary explicit, write

$$\hat{a}_m = 2 + 2 \cos k_m = |1 + e^{ik_m}|^2, \quad p = \sum_m \hat{a}_m.$$

The hat is not decoration: $\hat{a}_m$ is the *complement* of the $a_m = 4 \sin^2(k_m/2)$ used in the charge-odd shape basis (47) below, $\hat{a}_m + a_m = 4$, and that complementarity is the sum rule below.

Theorem 31 (Unsigned-incidence Bloch cubic). $p + q = 12$ *pointwise*, and the characteristic polynomial of $\mathcal{N}\mathcal{N}^\dagger$ is

$$\mu^3 - 2p\mu^2 + p^2\mu - 4\hat{a}_1\hat{a}_2\hat{a}_3 = \mu(\mu - p)^2 - 4\hat{a}_1\hat{a}_2\hat{a}_3,$$

with the unsigned comparison-adjacency eigenvalues $\lambda = \mu - 4$. The signed counterpart of the same computation is $\mu(\mu - q)^2$.

The algebraic closed form is complete rather than merely consistent at sampled momenta: the characteristic polynomial of the 81×81 and 192×192 finite plaquette adjacencies factorises coefficient for coefficient over $\mathbb{Z}$ into this cubic over the momentum grid.

Proposition 32 (Range and attainment). $\lambda \in [-4, 12]$ *exactly*. The top $\lambda = 12$ is attained only at Γ ($\hat{a}_1 = \hat{a}_2 = \hat{a}_3 = 4$). The bottom $\lambda = -4$ is attained on the whole of the three planes $k_j = \pi$, where the constant term vanishes, and only at R is it a triple root.

Proof. Write $f(\mu) = \mu(\mu - p)^2 - 4\hat{a}_1\hat{a}_2\hat{a}_3$. Then $f(0) = -4\hat{a}_1\hat{a}_2\hat{a}_3 \leq 0$ and $\mu(\mu - p)^2 < 0$ for $\mu < 0$, so no root lies below 0; and $f'(\mu) = (3\mu - p)(\mu - p) > 0$ for $\mu > p$, with $\hat{a}_1\hat{a}_2\hat{a}_3 \leq (p/3)^3$ by AM–GM and

$$f(16) \geq 16(16 - p)^2 - \frac{4p^3}{27} = \frac{4}{27}(12 - p)(48 - p)^2 \geq 0.$$

Hence $\mu \in [0, 16]$ and $\lambda = \mu - 4 \in [-4, 12]$. Equality at the top forces $p = 12$ and $\hat{a}_1\hat{a}_2\hat{a}_3 = 64$, which is Γ . Equality at the bottom is $f(0) = 0$, i.e. $\hat{a}_1\hat{a}_2\hat{a}_3 = 0$, i.e. some $k_j = \pi$; there $f(\mu) = \mu(\mu - p)^2$ and the root $\mu = 0$ is simple unless $p = 0$, which is R alone. $\square$

Remark 33. A flat unsigned-comparison branch would need $f(\mu_0) = 0$ identically in k , which the constant term already forbids. Its minimum is attained on three measure-zero planes, against the signed operator’s flat band. This contrast is an exact statement about two incidence constructions, not by itself a derivation of two physical charge sectors.

6.1 Conditional charge-even assembly

The shared-link coefficient at $r = 2$ is all-rank and analytic. Turning it into a charge-even band structure assumes the symbol identification stated above; the displayed second-order scalar additionally uses the supplied $SU(3)$ leakage/tower identification. The $r = 3$ values — $t_{3,+} = -6335/249696$ and everything computed from it — come from the same supplied domino ledger as Section 7, and are conditional on it in the same way.

Conditional on that identification, the scalar ledger is not a list of separate facts. Writing λ for the adjacency eigenvalue and $\text{tower}_{r,s}$ for the within-plaquette term in canonical u ,

$$E_s(\lambda, r) = \text{tower}_{r,s} + 12 \text{leak}_{r,s} + \lambda t_{r,s}, \quad (46)$$

with $\lambda \in \{-4, 8\}$ in the signed charge-odd sector (carrier and dispersive top), $\lambda \in \{12, -4\}$ in the unsigned comparison, and $\lambda = 0$ returning the on-site diagonals. This reproduces every registered diagonal and band value — nine of them — across both sectors and both orders. The tower term is the certified coupling conversion $4\Delta_{\pm}(3u/2)$ on its own sector’s series. The letter Δ is reused elsewhere with other subscripts — Δ_s for $s \in \{X, M, P, R\}$ are the fourth-order checkpoint differences of Section 9, Δ_C the gap between the two recorded off-axis coefficients, and Δ_L the finite-volume incidence separation of Section 11; only the sector subscript $\pm$ means this series. With that convention, (46) takes no unregistered input. The twelve is a count on this lattice, not a citation: the plaquette graph is 12-regular and two distinct faces share at most one link. The two towers are different objects and both are needed: the charge-odd one contributes $8/3 + u + u^2/2 + 7u^3/32$ and the charge-even one $8/3 - u + (13/20)u^2 + (101/200)u^3$, each of them $4\Delta(3u/2)$ on its own sector’s series, so $\text{tower}_{2,-} = 1/2$ against $\text{tower}_{2,+} = 13/20$ and $\text{tower}_{3,-} = 7/32$ against $\text{tower}_{3,+} = 101/200$.

Because the supplied ledger has $\text{leak}_{r,+} = t_{r,+}$ at both computed orders, the conditional charge-even assembly collapses to $\text{tower}_{r,+} + (12 + \lambda) t_{r,+}$, a one-parameter family in the charge-even tower. That equality is verified and unexplained: at third order the charge-odd leakage separates from the charge-odd hopping while the charge-even identity survives. It is recorded as a unifying candidate

with a falsifier — an order at which the two are computed independently and differ — not as a mechanism.

6.2 Where the rest frame is not blind

Conditional corollary 34 (Charge-even Γ splitting). If the unsigned symbol is the complete linked charge-even symbol, then at Γ its comparison spectrum is $\{12, 0, 0\}$, so the A_1^{++} level and the E^{++} doublet are distinct, and

$$E(A_1^{++}) - E(E^{++}) = 12 t_{r,+}$$

exactly: $-22/51$ at $r = 2$ and $-6335/20808$ at $r = 3$.

This is the structural counterpart of the charge-odd blindness recorded in Section 7. The signed spectrum at Γ is $\{-4, -4, -4\}$, one level, so no charge-odd rest-frame series can see the hopping; the unsigned comparison spectrum has two levels, so the charge-even rest frame could fix the coefficient under the symbol hypothesis. The A_1^{++} branch is isotropic at leading order without assuming it, with curvature $\frac{4}{3}|t_{r,+}|$ in all three directions — $22/459$ at second order and $6335/187272$ at third.

On a finite L^3 grid the conditional charge-even width is

$$W_{r,L}^{(+)} = |t_{r,+}| [12 - \lambda_{\min}^{(+)}(L)], \quad \lambda_{\min}^{(+)}(L) = -4 \iff L \text{ is even.}$$

Thus it equals $16|t_{r,+}|$ for even L and approaches that coefficientwise as $L \rightarrow \infty$; the charge-odd manifold width is $12|t_{r,-}|$ for even L . At any order at which the correction stays in the incidence algebra — which Proposition 38 warns is not automatic — the second-order ratio is $4(3N^2 - 5)/[3(N^2 - 4)]$, finite and greater than one at every rank $N \geq 3$, decreasing from $88/15$ at $N = 3$ towards 4. So at second order the charge-even band is the wider one at every rank under the charge-even symbol hypothesis, and that asymmetry is a channel statement rather than a geometric one.

7 $SU(3)$ through third order

The *third-order* content of this section is an externally verified computational input: the frozen upstream corpus reports a 251/251 cold exact replay for the retained chain, while this compact release checks only the assembled rationals. At second order the hopping coefficient t_3 is analytic, but the displayed flat scalar additionally uses the upstream $SU(3)$ identification $\text{leak}_2 = \ell_3$ and its within-plaquette tower. That tower contributes $8/3 + u + u^2/2 + 7u^3/32$.

Reported computational result 35 (Third-order ledger). The supplied ledgers report, in exact rationals,

$$b_3 = \frac{1975}{124848}, \quad \text{leak}_3 = -\frac{12331}{249696},$$

and the assembled flat scalar $d_3 = \frac{7}{32} + 12 \text{leak}_3 - 4b_3 = -\frac{109151}{249696}$.

Conditional corollary 36 (Third-order factorisation). If the lifter census is exhaustive and the linked ledger complete, then

$$H_{\text{eff},-}(k, u) = E_{\text{flat}}(u)I + \tau(u)B(k)B(k)^\dagger + O(u^4),$$

$$E_{\text{flat}}(u) = \frac{8}{3} + u + \frac{11}{306}u^2 - \frac{109151}{249696}u^3,$$

$$\tau(u) = \frac{5}{612}u^2 + \frac{1975}{124848}u^3.$$

The second-order diagonal is the $\lambda = 0$ value $\frac{1}{2} + 12\text{leak}_2 = 7/102$ with $\text{leak}_2 = -11/306$; subtracting $4t_3$ gives the flat scalar $11/306$ at $\lambda = -4$.

Remark 37 (What the external comparison can and cannot do). Hamer’s 1989 rest-frame series [14] agrees with E_{flat} at every shared order. It cannot test τ . The hopping enters the spectrum only through $\tau(u)q(k)$ and $q(0) = 0$, so every *charge-odd* Γ -point datum is blind to it. Under the unsigned-symbol hypothesis the charge-even rest frame is not, by Conditional result 34; the one-parameter family $(b_3 + \varepsilon, \text{leak}_3 + \varepsilon/3)$ leaves d_3 and therefore the entire comparison unchanged. What the agreement pins is the combination $12\text{leak}_3 - 4b_3$. The $\lambda = 8$ band top separates them, via (46). Conditional result 34 is the charge-*even* analogue and separates leak_+ from t_+ ; it says nothing about b_3 or leak_3 .

8 What the homology protects

Proposition 38 (Boundary-factorised rigidity). *If an order- r Bloch correction has the form $H_r(k) = a_r I + b_r S(k) + B(k)M_r(k)B(k)^\dagger$, then for every $k \neq 0$, $H_r(k)w(k) = (a_r - 4b_r)w(k)$: the correction shifts the carrier without dispersing it.*

Proof. $B^\dagger w = 0$ and $Sw = -4w$, both from Theorem 5. The middle term is annihilated whatever M_r is, which is the content: the protection belongs to the boundary operator and not to the dynamics. $\square$

In real space this is the orthogonality of the two cellular Laplacians, $L_2^\downarrow L_2^\uparrow = L_2^\uparrow L_2^\downarrow = 0$, so any polynomial in them acts on $\ker L_2^\downarrow \cap \ker L_2^\uparrow$ by its constant term. The proposition is not an unconditional all-orders statement: $\{BMB^\dagger\}$ is not a two-sided ideal in $\text{End}(C_2)$, and a general higher-order correction need not factor through links. Cubic symmetry alone permits carrier-projected fourth-order shapes built from

$$1, \quad q, \quad e_2, \quad \frac{4e_2}{q}, \quad \frac{e_3}{q}, \quad (47)$$

with $e_2 = \sum_{i < j} a_i a_j$, $e_3 = a_1 a_2 a_3$ and $a_i = 4 \sin^2(k_i/2)$. The $1/q$ is the carrier projection: $\|w\|^2 = q$ by Theorem 5, so the carrier-projected correction is $\varepsilon_4 = (w^\dagger H_4 w)/q$ and the numerator, not ε_4 , is the polynomial.

9 The fourth-order boundary

Two supplied fourth-order records agree on the axial datum and on the appearance of mobility, and disagree on one off-axis mixed-gradient coefficient C_{shp} . This edition does not adjudicate them, and no fourth-order band or bandwidth is used above. What it does is state precisely where the disagreement lives, because that geography is now checked and it is what a decisive calculation must target.

After subtracting the rest value, the recorded carrier ansatz is

$$\Delta_4(k) := \varepsilon_4(k) - \varepsilon_4(0) = Aq + Be_2 + C_{\text{shp}} \frac{4e_2}{q} + D \frac{e_3}{q}. \quad (48)$$

The continuous $k \rightarrow 0$ limit is used at $q = 0$. Denote the two recorded mixed-gradient values by C_{old} and C_{new} and set $\Delta_C = C_{\text{new}} - C_{\text{old}}$. A supplied all-rank cube ledger also reports

$$\alpha_N = \frac{640}{N(N^2 - 1)^3}, \quad (49)$$

but this formula and its identification with the complete axial coefficient remain conditional on that ledger. Sealed packages containing other fourth-order scalars use incompatible observable labels; none is promoted here to a canonical rest coefficient.

9.1 An exact theorem for the saved historical kernel

The upstream authority corpus supplies the hash-pinned 189-record historical kernel. Pulling its centered fourth-order operator back to cube amplitudes with $C = \partial_3$, define

$$Q_4 = C^\dagger H_4 C, \quad G = C^\dagger C, \quad Q_4 = Q_4 - s_4 G,$$

where $s_4 = \text{tr } H_4(\Gamma)/3$. If $L_i = 2I - T_i - T_i^{-1}$, then $G = \sum_i L_i$ and the centered coefficient is the generalized eigenvalue $Q_4 \phi = \lambda_4 G \phi$. The representative change $(Q_4, G) \sim (Q_4 + \delta G, G)$ is a scalar gauge only when the scalar anchor changes simultaneously; it cannot change an off-axis shape.

Reported computational result 39 (Exact saved-kernel Hodge pencil). For the supplied historical kernel,

$$q_{\text{old}}^{(4)} = -\frac{20721577909065127111}{7250590288602460800}, \quad \alpha_{\text{old}} = \frac{5}{12}, \quad \beta_{\text{old}} = \frac{17607806155349}{275331901291200},$$

and

$$Q_{4,\text{old}} = \frac{5}{48} \sum_i L_i^2 + \frac{17607806155349}{1101327605164800} \sum_{i < j} L_i L_j \succeq 0. \quad (50)$$

Its continuous-zone and even- L width is

$$W_{4,\text{old}} = \alpha_{\text{old}} + \beta_{\text{old}} = \frac{132329431693349}{275331901291200}.$$

These identities are exact for the saved kernel and its cold fixed-kernel checks. They do not establish that this kernel is the unique physical linked fourth-order operator.

The corresponding 25-point numerator stencil has

$$w_0 = \frac{9}{2}\alpha + 3\beta, \quad w_1 = -(\alpha + \beta), \quad w_2 = \frac{1}{4}\alpha, \quad w_d = \frac{1}{4}\beta,$$

and obeys $w_0 + 6w_1 + 6w_2 + 12w_d = 0$ exactly. For $\alpha, \beta > 0$, the generalized dispersion is bounded by $0 \leq \lambda_4 \leq \alpha + \beta$, with the continuous-zone extrema at Γ and R respectively; an odd- L grid does not contain R .

The two records do agree on the axial datum, but the corpus's own quoted tolerance does not cover that agreement: the α gap is 2.43×10^{-13} against a stated bound of 2.3×10^{-13} . The sealed core holds internally — α_{new} and $4A_{\text{new}}$ agree to 2×10^{-15} — which is consistent with the quoted bound being too tight, but two float comparisons do not discriminate that from a shared upstream slip. Recorded as a discrepancy rather than resolved, and the bound is not widened.

9.2 The obstruction is polynomial

Theorem 40 (Obstruction certificate). *The two recorded off-axis coefficients enter (47) only through the shape $4e_2/q$, so the two records' carrier numerators differ by $4\Delta_C e_2$ and their bands by $\Delta_C (4e_2/q)$. On any one-dimensional axial cut two of the three a_i vanish, so e_2 is the zero polynomial there and both differences vanish identically. Consequently no Γ -point or axial datum distinguishes the two records, at any precision. Off axis they do differ: $e_2(M) = 16$ and $e_2(R) = 48$, so the numerators differ by $64\Delta_C$ and $192\Delta_C$ and the bands by $8\Delta_C$ and $16\Delta_C$.*



Figure 3: Exact dispersion of the saved historical kernel in Reported result 39. The plot is a theorem about that hash-pinned kernel. Its identification with the complete physical linked $SU(3)$ fourth-order operator is disputed, so the curve is not used as a settled physical band.

Remark 41. The vanishing is polynomial, not numerical: this is non-identifiability, not a limitation of the data in hand. Re-anchoring cannot substitute, and neither record is preferred here. An off-axis contraction is the only decider.

On the axial cut itself the invariants collapse to $q = q_{\text{ax}}(k) = 4 \sin^2(k/2)$ with $e_2 = e_3 = 0$. Dividing the raw cube-boundary numerator $(\alpha_N/4)q_{\text{ax}}^2$ by $\|w\|^2 = q = q_{\text{ax}}$ leaves one power of q_{ax} with coefficient $\alpha_3/4 = 5/48$.

9.3 A second witness

Non-identifiability is better exhibited than argued from a difference between two records, since a difference could always be a computational error in one of them. The witness below is not itself a candidate for the physical coefficient. It is the balanced continuation obtained in the supplied source by setting an auxiliary bookkeeping parameter ϵ to zero; that source label carries no physical meaning here. Its sole role is to make the $4\Delta_{Ce_2}$ family concrete with an exactly known shift.

Proposition 42 (Explicit second witness). *There is a value C_{alt} with $C_{\text{alt}} - C_{\text{old}} = 25/1024$ exactly, appearing verbatim in the pinned source edition, whose induced band separations are 0 at X, 25/128 at M and 25/64 at R. It is identical to C_{old} on every retained datum and distinct from it off-axis.*

9.4 Where the coefficient lives

The four checkpoint deltas are the Γ -subtracted carrier corrections $\Delta_s = \varepsilon_4(k_s) - \varepsilon_4(0)$ at $k_X = (\pi, 0, 0)$, $k_M = (\pi, \pi, 0)$, $k_P = (\pi, \pi/2, 0)$ and $k_R = (\pi, \pi, \pi)$ — P is the off-axis, off-corner point that makes the system invertible, with $a = (4, 2, 0)$. Its algebraic role is explicit here; no additional

provenance is inferred for that checkpoint. In the shape basis they are

$$\begin{aligned}
\Delta_X &= 4A, \\
\Delta_M &= 8A + 16B + 8C_{\text{shp}}, \\
\Delta_P &= 6A + 8B + \frac{16}{3}C_{\text{shp}}, \\
\Delta_R &= 12A + 48B + 16C_{\text{shp}} + \frac{16}{3}D,
\end{aligned} \tag{51}$$

and they invert exactly, so the four shape coefficients are recoverable from four checkpoints. Δ_X contains A alone, which is why the axial cuts agree whatever the dispute does.

Four further measurements locate the coefficient and bound where a disagreement could live. The distinction matters and is easy to lose: showing where a coefficient's *value* is concentrated is not by itself showing where a *discrepancy* can sit. Only the third item below does the second thing.

1. **The shape fit's C row sums to zero**, so no error in the Γ anchoring can move C_{shp} at all; the recorded run's own re-anchoring moved it by 4.6×10^{-15} , as the zero row sum requires. A scalar shift is translation-local and changes nothing observable, which is why the two anchored scalars $q_{\text{band}}^{(4)}$ and $m_{\Gamma}^{(4)}$ are differently anchored coordinates and not rival estimates.
2. **The coefficient's value is concentrated, not spread.** Six of the 189 kernel records — the normal $(0, 0, 1)$ shell — carry 81% of the total weight, while the 120 rotation records carry 6% and are 255 times less sensitive per record. This says where the coefficient lives, not yet where the two records can differ.
3. **The agreed axial coefficient pins the normal sector.** All six normal records share one amplitude h , and it sets both shapes rigidly: $A_{\text{normal}} = -h$ while the raw normal amplitude is $2h$, which in the shape normalisation of (47) is $C_{\text{normal}} = h/2$, so

$$C_{\text{normal}} = -\frac{1}{2}A_{\text{normal}} \quad \text{exactly.}$$

(The factor of four between $2h$ and $h/2$ is the raw-to-shape conversion, not a discrepancy.) Granting $A = 5/48$ — which both records share — together with the block decomposition, which puts $5/48 + 4x$ on the normal block and $-4x$ on the mixed one, therefore fixes h and with it 81% of the coefficient; and everything free to move without disturbing A totals less than half the disputed gap. A rival kernel must differ *structurally* in the A -carrying sector; no redistribution reaches the gap.

4. **The two kernels agree almost everywhere.** Their supports are identical and 144 of 189 records ride one common scale; the divergent cross sector is a single uniform ratio. The dispute is three amplitudes.

Read together these four sharpen the target and prefer no side of it. They say that a rival kernel must differ structurally in the A -carrying sector rather than by redistribution. That is a constraint on candidates built on the recorded 189-record support — which is where the measurement was made; a kernel on a different support is not constrained by it at all. Neither record is promoted here and neither is withdrawn: the interval of *values* between them is not narrowed. What is narrowed is the class of *kernels* that can reach either end, and that constraint bears on both records alike.

9.5 What cannot decide it

A large sealed sweep over 609 marked clusters has been the standing candidate for settling this. It cannot: its engine assembles coefficients through a Γ -point scalar and emits no kernel block at

all, and by Theorem 40 Γ -point data are structurally incapable of constraining C_{shp} . This is a statement about the instrument, made by reading it, not a report of a failed run.

The instrument class that could eventually decide it now demonstrably exists: the sealed radius-two delivery of Section 5.4 emits a seed-space *matrix* at fifth order, not a scalar. What it lacks is off-axis momentum — a two-cube cluster has none — so it too cannot adjudicate. The requirement is now sharp on both counts: a decisive instrument must emit more than a Γ scalar *and* resolve an off-axis point such as P ; the marked-cluster engine fails the first condition, the radius-two delivery the second, and a target-blind calculation satisfying both is exactly the next step named in the conclusion.

9.6 The tier collapse, as two integers

The two shape coefficients B and D vanish in the recorded kernel even though the two-hop enumeration generates both, and the degree count does *not* forbid them. The carrier numerator reaches a -degree three, because the projection contributes one further power of a : for any vector whose components have squared modulus a_i — the carrier is one, in whichever of the two equivalent labellings — the quadratic form $\text{diag}(a_i^2)$ evaluates to $\sum_i a_i^3 = q^3 - 3qe_2 + 3e_3$, an explicit witness carrying a nonzero e_3 . An earlier degree-bound mechanism proposed for the collapse was therefore retracted. What is checked is finer: every degree-three block amplitude is an integer multiple of one raw record weight x , and

$$B = 0 : (+1, +1, -2) \cdot x, \quad D = 0 : (-3, +6, -3) \cdot x$$

over the displacement shells. The kernel does populate the forbidden tier; the dynamics cancels it with small integers. Any mechanism must reproduce those two vectors.

9.7 A near- Γ statement that does not adjudicate

One quantitative fourth-order statement can be made without adjudicating anything, because it depends on the kernel only through $\sqrt{W_4}$. Here $W_4 = \sup_k (\varepsilon_4(k) - \varepsilon_4(0))$ is the zone spread of the carrier-projected fourth-order correction about Γ , and $\theta \in (0, 1)$ is the fraction of the second-order isolation gap the fourth-order spread is allowed to occupy; we take $\theta = 1/2$ throughout.

What the criterion needs is a lower bound on q over the region $|k| \geq r$, and there it is sharp: writing $q = \sum_m h(k_m^2)$ with $h(t) = 2(1 - \cos \sqrt{t})$, the identity $(s \cos s - \sin s)' = -s \sin s \leq 0$ makes h concave and increasing on $[0, \pi^2]$, so the minimum over the simplex sits at a vertex and

$$\min_{|k| \geq r} q(k) = 4 \sin^2(r/2), \tag{52}$$

attained on a coordinate axis. With $\tau(u) \geq t_3 u^2$ for $u > 0$, the fourth-order shape terms are therefore dominated by the second-order dispersion outside a ball of radius $2 \arcsin(u \sqrt{W_4}/(\theta t_3))$, which is Ku to leading order with $K = \sqrt{W_4}/(\theta t_3)$: the two records give $K = 10.85$ and $K = 15.06$ — floating point, since one of the two bandwidths exists only as a float — and the larger bounds both, so the exclusion radius can be stated for either record without choosing between them.

Jordan's inequality $q(k) \geq (4/\pi^2)|k|^2$ also holds on the whole zone and is tight at $k = 0$ and at the corner, but the corner is not where this bound is used; on the axis it is loose by $\pi^2/4$, and the constants it gives, 17.04 and 23.66, are exactly $\pi/2$ times the sharp ones.

It is not a bound on the open coefficient: a value of C_{shp} outside the two records carries its own W_4 , and nothing here confines the answer to the interval the records span. The statement is non-vacuous only for $u < 0.133$ at $\theta = 1/2$, and that threshold is the same for both constants —

the sharp radius reaches π when $u\sqrt{W_4/(\theta t_3)} = 1$, which is exactly where $Ku = \pi$ for the Jordan constant, because the two bounds agree at the zone corner. The excluded zone fraction grows as u^3 .

10 External comparisons

External results count here because they are *independent* — produced without knowledge of this program — not because they are published. We use them only within the scope stated below; all normalisation conversions and coefficient matches are calculations of this paper.

Hamer 1989, both sectors. Let a_n denote the coefficient of x^n in Hamer’s tabulated dimensionless mass series. The 1^{+-} rest-frame series agrees with E_{flat} at every shared order: $16/3 \div 2 = 8/3$ exactly, $a_1 = +1$, and $2a_2, 4a_3$ matching $11/306$ and d_3 to 9.9×10^{-14} and 1.1×10^{-12} . Conditional on the unsigned-symbol identification, the 0^{++} series also matches the charge-even A_1^{++} Γ -point coefficients at x^2 and x^3 , with the sign structure ($a_1 = -1$ in the scalar channel where the carrier’s is $+1$) reproduced; the E^{++} doublet is not constrained by it, which is why Conditional result 34 uses the splitting rather than the level. The conversion $m_n = 2^{n-1}a_n$ is $x = 2u$ together with Hamer’s normalisation $ma = (g^2/2)M$ — the 2^n from the substitution, the half from the normalisation — exactly, with no 4^r ambiguity anywhere. Hamer states in print that his x^3 and x^4 terms disagree with the earlier series of [4]. We report Hamer’s comparison and do not independently adjudicate the earlier coefficient table.

The string tension. Let t_n denote the coefficient sequence in the transcribed Kogut–Pearson–Shigemitsu $SU(3)$ table [7] — read here from the 1980 preprint, the journal pages themselves being unread — and let σ_n denote the corresponding ledger coefficient after the displayed convention conversion. The transcription satisfies exactly $2^{n-1}t_n = \sigma_n$ for $n = 2, \dots, 5$, in the physical sign convention $\sigma_n^{\text{phys}} = (-1)^n \sigma_n^{\text{raw}}$. The same series reproduces E_{flat} term by term through u^3 , including the discriminating u^3 coefficient.

Euclidean series. Historical Euclidean strong-coupling series by Münster, Seo and Smit [6, 8, 9, 10, 11] provide context only. Two printings of the same Münster/Seo series — the second erratum to [6], which prints the definitive corrected table and credits Ukawa and Seo [8] with locating the error, and Smit’s Table 1 scalar column [9], which credits them — agree rational for rational through u^8 . The accompanying corpus also contains an apparent eighth-order discrepancy between the 1982 erratum and the 1985 table. Because every original table page has not been independently re-transcribed for this paper, neither the apparent equality nor the discrepancy is used as evidence. No Hamiltonian coefficient is inferred from these Euclidean records.

The overlap obstruction. Schierholz [13] measured in 1988 that at $\beta = 5.9$ — in the Euclidean scaling region, not at strong coupling — a bare local operator carries only a couple of percent of the state it interpolates, and gave the scaling a^5 for a local dimension-four operator. Together with Brandstaeter, Kronfeld and Schierholz [15] these document the degradation of local-operator projection and its critical signal-to-noise scaling as the lattice spacing is reduced. Nothing quantitative is transferred to the present strong-coupling finite-order setting. Read only as a warning, these results motivate posing $\mathbb{G}18$ in a smeared basis; they supply no bound and do not identify the carrier with a physical glueball.

Cross-regime discipline. Two indexed papers declare an explicit scope firewall and neither supplies a value here: Davoudi–Raychowdhury–Shaw [23], whose Hilbert-space cost comparison is locked to 1+1 dimensions with matter, and Kadam–Raychowdhury–Stryker [24]. The rule is enforced by the literature validator rather than recorded in prose.

11 Scope

This work does not establish a four-dimensional infinite-volume Yang–Mills mass gap, a continuum glueball mass, convergence of the strong-coupling series to the continuum, regulator independence of the first dispersive order, the full $SU(3)$ fourth-order rest coefficient or off-axis shape, or an identification of any level here with a physical glueball. At Γ the three face orbitals transform as an axial vector of the cubic group, which with parity even and charge conjugation odd gives the retained-space label T_1^{+-} ; this is an operator-space assignment, not a spectral one.

The carrier is macroscopically degenerate at finite L . Under the third-order factorisation input, the minimum internal separation of one sheet from the incidence sheets is $\Delta_L(u) = 4\tau(u) \sin^2(\pi/L) + O(u^4)$, which falls as L^{-2} ; it cannot provide a volume-uniform contour for that single sheet over the entire Brillouin zone. This statement must not be conflated with external isolation of the complete electric shell. Theorem 1 gives the latter a volume-independent free margin $C_F/2$, while Proposition 10 gives an L -independent internal coefficient at every fixed nonzero nested momentum.

Three crossings are named explicitly, and none of them is made here.

Proposition 43 (Exact reduction of the G17 free-energy target). *For the Euclidean $SU(3)$ Wilson plaquette density $W_p = 1 - \frac{1}{3} \text{Re Tr } U_p$, let*

$$Z_\Lambda(b) = \int \exp\left(-\sum_p b_p W_p\right) d\mu, \quad b_p(s) = \begin{cases} s, & p \in \Gamma, \\ \beta, & p \notin \Gamma, \end{cases}$$

and write $Z_{\beta,\alpha,\Gamma} = Z_\Lambda(b(\alpha\beta))$ and $Z_\beta = Z_\Lambda(b(\beta))$. Then

$$\log \frac{Z_{\beta,\alpha,\Gamma}}{Z_\beta} = \int_{\alpha\beta}^{\beta} \sum_{p \in \Gamma} \langle W_p \rangle_{s,\Gamma} ds. \quad (53)$$

Consequently the single local estimate

$$\sup_{\Lambda, \Gamma, p \in \Gamma, s \in [\alpha\beta, \beta]} s \langle W_p \rangle_{s,\Gamma} \leq C \quad (54)$$

would imply the desired animal-uniform bound with

$$\frac{Z_{\beta,\alpha,\Gamma}}{Z_\beta} \leq K_\alpha^{|\Gamma|}, \quad \log K_\alpha \leq C \log(1/\alpha). \quad (55)$$

Proof. Differentiating $\log Z_\Lambda(b(s))$ gives $-\sum_{p \in \Gamma} \langle W_p \rangle_{s,\Gamma}$ and integration from β down to $\alpha\beta$ proves (53). Under (54), each summand is at most C/s , which proves (55). $\square$

The elementary pointwise estimate $0 \leq W_p \leq 3/2$ gives only $\log K_\alpha \leq \frac{3}{2}(1-\alpha)\beta$; its growth in β is precisely why the uniform C/s estimate remains a substantive open problem.

- **Finite lattice to infinite volume** requires an inhomogeneous Wilson free-energy bound with a useful $\log K_\alpha$ and a source-radius reduction. Proposition 43 turns the first item into the explicit one-plaquette estimate (54), but does not prove it; the source-radius theorem is

also still unstated. Both are carried in the accompanying gap register as G17. The Hamiltonian Riesz island of Theorem 2 belongs to small u and does not prove this Euclidean weak-coupling bound.

- **Protected level to particle** requires a volume-uniform overlap of a smeared, dressed T_1^{+-} operator built on the carrier with the transfer-matrix spectrum, and multi-plaquette survival of the protected level (G18). The published overlap measurement of Section 10 is why this theorem must live in the smeared basis. This edition now settles the free external gap, the complete finite-volume interacting Riesz island at existentially small $|u|$, the fixed- k carrier, and the exact price of combining locality with sharp momentum. It does not prove a volume-uniform Wilson-source frame, infinite-volume totality, or the transfer map.
- **Lattice to continuum** (G19) is untouched. Small u is strong coupling; the continuum is the opposite end of the axis.

All three crossings are required and none is completed here. The new window and Riesz-island theorems remove the free-spectrum obstruction and its small- u finite-volume persistence subproblem from G18, but they do not promote the protected lattice cluster to an infinite-volume physical particle.

The two former charge-odd completeness premises are now theorems. Theorem 1 proves more than shell completeness: below $\frac{5}{2}C_F$ the full trivial-flux electric spectrum is exactly $\{0, 2C_F\}$ for every $N \geq 3$ and $L \geq 3$, and the range of P_- is exactly the charge-odd one-plaquette span. Lemma 16 proves that the only off-diagonal second-order routes are the four adjacent shared-link channels: disjoint exchange routes are orthogonal and the same-face route cancels. The explicit projector matrix element (36) then fixes their signs and normalisations. These statements are analytic results; the finite-volume $L = 3, 4, 5$ enumerations in the verifier are regressions, not substitutes for their proofs. They are new in this edition and have not yet faced external referees; a reader auditing this paper should start there.

One earlier physical identification remains conditional.

That the complete linked charge-even Bloch symbol is the unsigned incidence $\mathcal{N}(k)$ — the signed incidence with the boundary orientations dropped — is not derived. The cubic and range in Section 6 are exact algebraic properties of the defined comparison operator; only their physical charge-even interpretation, bandwidth and Γ -splitting rest on this additional hypothesis.

12 Evidence, authority, and reproducibility

The accompanying dependency-free program `verify_publication_core.py` uses exact integers and `Fraction` arithmetic. Its thirty-eight checks are the coefficient ledger and Weingarten sums, signed incidence and torus homology, closed-cube shell, two-cube connected geometry and reported channel arithmetic, the unsigned cubic, the saved historical Hodge pencil, the fourth-order obstruction, the coupling convention, the planar closed form (38), and the exhaustive shared-link geometry, cycle censuses and torus ranks through $L = 5$ — every value a `Fraction` and no float constructed anywhere in the file.

It does *not* cover the arithmetic new in this edition: the all-rank Casimir shelf of Theorem 1, the $SU(3)$ additive-spectrum and Riesz-contour arithmetic of Theorem 2, the cycle censuses those proofs delegate to, the fixed-momentum carrier, the harmonic representatives and the sharp zone minimum. Those are checked in the accompanying repository instead, and each carries its own printed marker; the boundary is stated rather than blurred, because a standalone verifier that is described as covering more than it runs is worse than one with a stated edge. A second program,

`verify_radius2_report.py`, holds everything the sealed radius-two artifact of Section 5.4 can support: it is also standard-library only (the payload is read with `zipfile` and `struct`, and the 11×11 block is diagonalised by cyclic Jacobi), but it works in floats at stated tolerances, and the two files are kept separate precisely so that the exact one stays float-free. Neither reconstructs the finite-precision Clebsch–Gordan inputs, replays the third-order microscopic contractions, or chooses C_{shp} .

Every sentence of this paper that a machine check certifies carries an inline `\chk` marker in the source; the markers are extracted verbatim into Appendix G, so the coverage claim is itself auditable rather than asserted. In the accompanying repository `workhouse verify -only '<name>'` reruns any one of them alone, with its numbers, in about a second.

That device is itself checked. A registered invariant reads every pinned edition in `paper/` — this one included — resolves every printed label against the live registry, and fails if one is renamed, deleted, or left red. The guard is the reason this edition’s labels can be trusted at all: the same invariant read only the earliest pinned edition until this one landed, and nineteen labels that had drifted in the intervening editions passed unnoticed until it was made to read them all.

A few displayed objects deliberately carry no marker, because none is a result of this program: the dimension and Casimir table (25), which is standard $SU(N)$ representation theory; the boundary formulas of Appendix E, which are definitions — what is checked there is that they compose to zero; and Table 1, which is provenance about a commit rather than a claim about the physics.

Layer	Count
T0 — Lean 4 theorems, no <code>sorry</code>	37
T1 — exact re-derivations	191
T2 — numerical, within a stated tolerance	48
Repository tests	484

Table 1: Counters at the pinned commit. The T0 row counts theorem declarations with no `sorry`, which is a scrape of the tree; `make lean` is what compiles them, and the standard the tree is held to is that every axiom footprint lies in `{propext, Classical.choice, Quot.sound}`. That compilation was last run elsewhere: the environment this edition was typeset in carries no Lean toolchain, and the row is not asserted as remeasured here. The T1 and T2 rows are `make verify` and the test row is `make check`. They are provenance rather than independent evidence; the argument is above, and the checks are the checks.

A smaller layer is proof-checked rather than merely re-derived. The Lean development compiles the rank law with denominators cleared, $t_3 = 5/612$ and $t_2 = 0$, the third-order ledger identity, the axial law at every exceptional rank, the cycle count $(L^3 - 1) + 3 = L^3 + 2$, the four checkpoint-extraction formulas and the deltas they invert, and four statements about the charge-even cubic including the identity that its middle coefficient is p^2 . It is exact-rational and polynomial algebra, plus the trigonometric checkpoint evaluations over $\mathbb{R}$ that supply those deltas — the $\mathbb{R}$ section exists because the extraction theorems formalised only half a statement without it. There is no perturbative derivation and no operator theory. An earlier edition added that “nothing that bears on C_{shp} appears there”. That was false, and false about the one genuinely open item here. The tree carries $C_{\text{shp}}^{\text{old}}$ as a rational constant, proves it from the pencil coefficients ($C = (\beta - 2\alpha_3)/16$), and proof-checks the four checkpoint-extraction formulas of Section 9 — which is precisely the C_{shp} algebra. What remains true is the weaker and duller statement: nothing in the Lean tree *decides* between the two recorded values, and by Theorem 40 nothing stated on Γ or an axis could.

12.1 Claim-level evidence map

Claim	Status	Evidence in this edition
Signed incidence spectrum and $L^3 + 2$ carrier	Proven	Analytic integer topology; exact replay
Adjacent all-rank coefficient t_N	Proven locally	Weingarten derivation; exact replay
Global order- u^2 torus assembly	Proven	Retained-shell theorem; process-completeness lemma; explicit projector matrix element
Uniform first electric spectral window	Proven	Spin-network support census; all-rank Casimir shelf; trivial one-form flux
$SU(3)$ all-orders finite-volume Riesz island	Proven, existential do-main	Yarotsky enclosure; exact additive spectrum; periodic bounded-degree transport
Two-cube six-channel census	Reported numerical	Rational reconstructions; shared amplitude lineage; assembly test
Unsigned-incidence cubic and range	Proven algebraically	Exact characteristic polynomial and finite-grid replay
Physical charge-even identification	Conditional	Complete linked symbol, including vacuum route, not derived
$SU(3)$ factorisation through u^3	Upstream cold-exact	251/251 exact gates reported by the frozen authority; raw replay not bundled here
Historical fourth-order pencil	Exact for saved kernel	Hash-pinned payload and cold fixed-kernel checks
Complete physical fourth-order kernel	Disputed/open	Historical and August candidates differ off axis
Radius-two rational sub-block and C_4/C_5 increments	Reported numerical	Sealed hash-pinned artifact; stdlib float verifier at stated tolerances; detached replay of pinned parents; no odd Q_3 enclosure
Exact G17 reduction	Proven reduction; estimate open	Differentiated inhomogeneous partition function; missing uniform one-plaquette bound
Infinite-volume particle or continuum gap	Open	No thermodynamic Wilson frame, transfer identification, or continuum bridge

Table 2: Truth status and evidence level are separate. “Exact for a saved kernel” does not establish that kernel’s physical identification.

12.2 Upstream corpus and downstream verifier

The organised ALL THEORY archive was consulted as the upstream scientific corpus. Its four-document v4.3 authority stack is frozen by byte count and SHA-256 in `export/MAN_GOV_export_manifest.csv`; all four entries were rehashed for this edition and matched. The downstream verifier is [WORKHOUSE](#). Document flow is one-way—frozen corpus to hash manifest to verifier—while verifier findings return to review and status records. The later two-cube results are not in the August 20 corpus and are therefore kept as later WORKHOUSE-local evidence.

The current hostile-scope audit branch is pinned at `ff9a5976a5b03cd763e3e50daf4326138770460d`. On Windows, its registry was rerun locally: 229/230 registered checks passed. The sole failure was the known path-separator comparison for a restored payload; the payload hash and all paper-relevant arithmetic checks passed. This is a platform-specific local audit result, not a claim

that the complete test suite or Lean tree was rerun. The source-level `\chk` labels remain in this TeX file but are suppressed typographically; this revised edition is not claimed to be covered by the repository’s older manuscript-label guard.

The authority files are identified by the following SHA-256 values:

- master v4.3: 1f8295c3a797aeb854b4b0ca82485bd7f4a149617777d2a43fc147a4a53c26dd;
- detailed formulas v3.1: 924d5b26f865e33f530d166735e847b89f60df7fe363995e2788ab24e1ff4254;
- consolidation guide v4.3: 60b8a93c8149049f50ed3ce8e0cfd1b379723b6765f201395a2fda6c4a7e7aa7;
- canonical source manifest v4.3: 730bc99ed8d749de0dffbabfda85cd4512d8f7748782620b694759a03bacde05.

12.3 Data and code availability

This release directory contains the TeX source, generated vector figures, figure generator, exact verifier, build log, rendered-page audit, and a SHA-256 manifest. The historical kernel payload used in Reported result 39 has SHA-256 635d40fa8a5d7da841fd30f36185eb96f14ec4c040678ddd8fb010379afb2900. The pinned ymcirc $d = 3$, $\mathcal{B} = 4$ data blob is identified in Ref. [29] and has SHA-256 4cad9538f2c052488556362e0cf6e4e63363e7f5801597aeab9013ef945c59dc. The exact stranded-flux audit source and result have SHA-256 a42800a91809a77bbd1d339ac604ea80a234eed305377bbabbbf890fee98fde55 and d215fa1e9cd0bde0410eec16c687182c30bc62fa3f28c1445fd81e93888e1d60. The sealed radius-two spectrum of Section 5.4 and its certificate are bundled, with SHA-256 52a17f87d3990b234635828a0c92f73f781e9a3abec8bb3c127e990b7a717e6c and b9b14d28438585fc4f7d88f96497e9c6e01d0532a8402ffd6db17258a1638c0e; `verify_radius2_report.py` refuses to run against any other bytes. The upstream release governs the wider supplement through `two_cube_cutoff_free_radius2_finite_u_spectrum_manifest.json`, which pins twelve members — derivation note, builder, hostile tests, six sealed-parent files, result NPZ, certificate and detached verifier — of which only the two named above travelled with this edition. The two-cube build and its finite-precision Clebsch–Gordan payloads are not reproduced by the compact verifier; the exact statements that depend on them are labelled reported. A durable tag or archival DOI is still needed before journal submission.

Machine checks certify algebra relative to their definitions and inputs. They do not replace derivations, promote numerical contractions to symbolic ones, or certify an uncompleted off-axis contraction.

13 Conclusion

The strongest exact coefficient result is a finite-order operator theorem built on a volume-uniform free spectral window: for every $N \geq 3$ and periodic $L \geq 3$, the trivial-flux electric spectrum below $\frac{5}{2}C_F$ is exactly $\{0, 2C_F\}$, the retained-shell theorem identifies the full charge-odd $2C_F$ eigenspace, the process-completeness lemma exhausts its second-order off-diagonal routes, and the four shared-link fusion projectors give the closed coefficient t_N with the oriented cellular boundary fixing its sign and geometry. This assembles *without* a completeness hypothesis into a homological flat carrier of dimension $L^3 + 2$: the flat sector is not a fine-tuned interference band but the cycle space of the lattice two-complex, degenerate because $\partial_2\partial_3 = 0$. Those proofs are the newest and least-reviewed material here, and they deserve the reader’s hardest look.

The apparent L^{-2} obstruction is now localised: it is the minimum internal sheet separation near Γ , whereas the free external shell margin is at least $C_F/2$ and the coefficient at every fixed nonzero nested momentum is volume-independent. The $SU(3)$ two-cube calculation supplies a nontrivial operator-level assembly check, including channel-by-channel geometry and orbit localisation, but shares its amplitude lineage and remains numerical at the contraction layer.

For $SU(3)$ the analysis goes one step beyond finite order. Yarotsky's local spectral enclosure, together with the exact additive Casimir arithmetic and a periodic bounded-degree transport, gives an all-orders rank- $3L^3$ Riesz island for the complete neutral charge-odd cluster whenever $27|u| < c_1$ and $27c_2|u| < 1/33$. This is a volume-uniform finite-volume statement with existential constants. It neither isolates one internal sheet nor proves a source-carrying thermodynamic particle. A later detached-replayed radius-two extension adds nested finite- u Ritz spectra and leading incremental order- u^4 and order- u^5 effects; because the odd exterior action is absent, it supplies neither a matched truncation error nor a complete fourth- or fifth-order physical coefficient.

The unsigned-incidence cubic is exact as an algebraic comparison; its physical charge-even identification is not assumed silently. At fourth order the saved historical kernel has an exact positive generalized Hodge pencil and stencil, yet the later candidate differs in an off-axis coefficient that no scalar reanchoring, Γ datum, or axial datum can determine. The decisive next object is therefore one target-blind, hash-sealed calculation that returns the rest scalar and off-axis shape from the same kernel. Independently, promoting the finite-volume Riesz island to a thermodynamic, source-carrying multiplet with a volume-uniform Wilson frame and a controlled transfer map is now the precise fixed-spacing **G18** target, and Proposition 43 isolates **G17**'s first missing estimate without pretending to prove it. Until those steps are supplied, the operator and Riesz-island theorems are substantial, but the physical fourth-order kernel, infinite-volume particle interpretation, and continuum limit remain open.

A Channel weights in closed form

Substituting (25) into (32),

$$\begin{aligned} w_{\mathbf{1}} &= -\frac{2}{N(N^2-1)}, & w_{\Lambda^2 F} &= -\frac{N-1}{(N+1)(2N-3)}, \\ w_{\text{Adj}} &= -\frac{2(N^2-1)}{N(2N^2-1)}, & w_{\text{Sym}^2 F} &= -\frac{N+1}{(N-1)(2N+3)}. \end{aligned}$$

Their family sums are (33) and (34), and the difference has numerator $2N(N^2-4)$, which is the $SU(2)$ zero and the positivity for $N \geq 3$ at once. At $N=3$: $w_{\mathbf{1}} = -1/12$, $w_{\text{Adj}} = -16/51$, $w_{\Lambda^2 F} = -1/6$, $w_{\text{Sym}^2 F} = -2/9$, which are the four values Reported result 21 matches against (42).

B The Weingarten index sums

With $\text{Wg}(e) = 1/(N^2-1)$ and $\text{Wg}((12)) = -1/(N(N^2-1))$, and σ, τ running over S_2 ,

$$\begin{aligned} &\langle U_{i_1 j_1} U_{i_2 j_2} \bar{U}_{i'_1 j'_1} \bar{U}_{i'_2 j'_2} \rangle \\ &= \sum_{\sigma, \tau} \text{Wg}(\sigma\tau^{-1}) \prod_a \delta_{i_a i'_{\sigma(a)}} \delta_{j_a j'_{\tau(a)}}. \end{aligned}$$

Summing the delta products over all N^4 index quadruples,

$$M_{\text{direct}} = \text{Wg}(e)(N^4 + N^2) + 2 \text{Wg}((12))N^3 = N^2,$$

$$M_{\text{cross}} = 2 \text{Wg}(e)N^3 + \text{Wg}((12))(N^4 + N^2) = N.$$

Both are confirmed by explicit summation at $N = 3, 4, 5$. The related fourth moment $\int |U_{11}|^4 dU = 2/[N(N+1)]$ is $1/6$ at $N = 3$ and in particular nonzero, which is what refuted a claimed Haar vanishing in the balanced $(2, 2)$ sector.

C The $SU(3)$ coefficient ledger

Order	Quantity	Value
u^0	electric plaquette	$8/3$
u^1	scalar	1
u^2	$BB^\dagger$ coefficient	$5/612$
u^2	flat scalar	$11/306$
u^3	$BB^\dagger$ coefficient	$1975/124848$
u^3	per-neighbour leakage	$-12331/249696$
u^3	flat scalar	$-109151/249696$

The u^1 row is $SU(3)$ -only rather than an $SU(3)$ specialisation: it is $-\langle p, -|V|p, -\rangle = 1$, which needs the ϵ channel and vanishes for $N \geq 4$. Of the rest, the u^0 row and the u^2 $BB^\dagger$ coefficient specialise all-rank statements ($2C_F$ and t_N); the u^2 flat scalar inherits the $SU(3)$ tower term through (46); and the three u^3 rows are $SU(3)$ -only and conditional on the supplied domino ledger, exactly as Section 7 says.

D Unsigned comparison and conditional charge-even ledger

Conditional on identifying the unsigned-incidence comparison with the complete linked charge-even symbol, the same convention, with $\lambda \in \{12, -4\}$ and (46), gives:

Order	Quantity	Value
u^2	hopping $t_{2,+} = \ell_3$	$-11/306$
u^2	on-site diagonal	$223/1020$
u^2	band bottom ($\lambda = 12$)	$-217/1020$
u^2	band top ($\lambda = -4$)	$1109/3060$
u^2	bandwidth $16 t_+ $	$88/153$
u^3	hopping $t_{3,+}$	$-6335/249696$
u^3	on-site diagonal ($\lambda = 0$)	$52163/260100$
u^3	band bottom ($\lambda = 12$)	$-54049/520200$
u^3	band top ($\lambda = -4$)	$471353/1560600$

Because $t_{r,+} < 0$ at both orders, the $\lambda = -4$ endpoint is the band *maximum* in this sector. A certificate key naming it **bandmin** at both orders is a naming error against its own arithmetic, recorded rather than silently corrected.

E Finite-volume chain calculation

For $i < j$,

$$\partial_2[x; i, j] = [x + e_i; j] - [x; j] - [x + e_j; i] + [x; i],$$

$$\begin{aligned}\partial_3[x; 1, 2, 3] &= [x + e_1; 2, 3] - [x; 2, 3] \\ &\quad - [x + e_2; 1, 3] + [x; 1, 3] \\ &\quad + [x + e_3; 1, 2] - [x; 1, 2].\end{aligned}$$

Substituting the second into the first cancels every edge pairwise, so $\partial_2\partial_3 = 0$ over $\mathbb{Z}$. On the torus $\ker \partial_3$ is generated by the sum of all oriented cubes, so $\text{rank } \partial_3 = L^3 - 1$.

F The two-cube geometry

The eleven oriented plaquettes of the $(3, 2, 2)$ prism, by base vertex and ordered plane axes, are

$$\begin{aligned}(000; 12), (000; 13), (000; 23), (001; 12), \\ (010; 13), (100; 12), (100; 13), (100; 23), \\ (101; 12), (110; 13), (200; 23),\end{aligned}$$

with left, right and shared-face index sets $I_L = (0, 1, 2, 3, 4, 7)$, $I_R = (5, 6, 7, 8, 9, 10)$, $I_F = (7)$. Building $G = BB^\top$ from the oriented boundaries and applying the geometric Möbius fold (40) leaves exactly four nonzero off-diagonal pairs, $(0, 5), (1, 6), (3, 8), (4, 9)$, each equal to -1 : four disjoint 2×2 blocks and three isolated directions, which is why the spectra in Section 5 are closed form once the diagonals below are supplied. The two connected diagonals are

$$\begin{aligned}D_{\mathcal{B}=6} &= \frac{1}{612} \text{diag}(-2317, -2317, -2295, -2317, -2317, \\ &\quad -2317, -2317, 0, -2317, -2317, -2295), \\ D_{\mathcal{B}=4} &= \text{diag}\left(-\frac{7}{4}, -\frac{7}{4}, -\frac{15}{4}, -\frac{7}{4}, -\frac{7}{4}, \right. \\ &\quad \left. -\frac{7}{4}, -\frac{7}{4}, 0, -\frac{7}{4}, -\frac{7}{4}, -\frac{15}{4}\right),\end{aligned}$$

so both spectra are one line of arithmetic: a pair block contributes $D_{ii} \mp c$ twice and each isolated direction its own D_{ii} . The $\mathcal{B} = 4$ comparator on the same geometry gives 0 once, $-\frac{15}{4}$ twice, $-\frac{11}{6}$ four times, and $-\frac{5}{3}$ four times: the paired levels sit *above* the isolated $-15/4$ pair, where at $\mathcal{B} = 6$ they sit below it. The three isolated levels $\{-\frac{15}{4}, -\frac{15}{4}, 0\}$ are identical in the two truncations, so the entire effect of the cutoff lives in the four cross-cell blocks. That reversal is not the coefficient's sign by itself — within a 2×2 block the pair splits as $D \mp c$, which is invariant under $c \rightarrow -c$ — but the truncation's effect on the connected diagonal as well. What the sign of c alone does reverse is the ordering on a *single* cube, where the kernel is $\alpha I + tG$ with one diagonal and the spectrum of G fixed; that is the shell reversal the one-cube reconstructions report.

G Machine-check coverage

The paper's discipline is that a machine check, not a sentence, is the unit of evidence. This appendix makes that auditable: every inline `\chk` marker in the source — each one naming the check that certifies the sentence it sits beside — is extracted verbatim by `make_coverage.py` and printed here. A reader can therefore see exactly which statements are covered, and, by their absence from this list, which are not: the remaining load-bearing hypothesis of Section 11 — the charge-even identification — appears nowhere below, and that is the point.

The source carries 126 inline markers. Each line below is one marker, verbatim, under the section that carries it.

Introduction

- the discrete index theorem gives 0 on the signed face-edge operator, bounding nothing

Hamiltonian and projected sector

- the printed towers are canonical-u: $4\Delta(3u/2)$ reproduces them verbatim
- the 4^*r rescaling breaks the bridge: order 2 off by 16, order 3 by 64
- at $N = 2$ the C-odd hopping vanishes and the C-even one does not
- every nontrivial irrep clears $5/4 C_F$, except the fundamental pair
- the torus graph is simple, and the only short cycles are the $L = 3$ winding triangles
- every 4-cycle is an elementary face, except exactly the $3L^2$ straight winding loops at $L = 4$
- the $SU(3)$ additive electric spectrum leaves 14, 16, 17 around the shell

Oriented incidence and the exact carrier

- $B^\dagger = q I - d \text{conj}(d)^T$ for the curl incidence
- $\dim Z_2 = L^3 + 2$ by rank, not by re-arranging the formula
- cube boundaries and three wrapping sheets $\text{SPAN } Z_2$
- the L^3+2 count is chain-level, not the Bloch convention
- the Bloch and chain routes to the carrier agree
- the sheets average to harmonic representatives, and the Gamma block IS b_2
- the Fourier sum of cube boundaries IS the carrier, and a local seed costs L^3

The wrapping sheets are cycles, not harmonic — and their harmonic representatives

- FINDING: the wrapping sheets are cycles but NOT harmonic
- the sheets average to harmonic representatives, and the Gamma block IS b_2

Exact adjacent shared-link dynamics and global assembly

- each channel gap is $C_F + C_R/2$, and the weights sum to one

The channel weights are a theorem

- the shared-link weights are Weingarten, not an isotropy assumption
- the Weingarten route is independent of the corpus
- the published dimension-ratio matrix element is this registry's weight formula
- the fourth moment integral $|U_{11}|^4 = 1/6$ at $N = 3$

The off-diagonal process list is complete

- the plaquette graph is 12-regular and two faces share at most one link
- connected first order vanishes by inclusion-exclusion, not by cancellation of numbers
- the four channel weights follow from dimension and Casimir
- A_N and B_N are the channel sums, not transcriptions

The shared-link law

- $t_N = B_N - A_N$
- $t_N > 0$ for $N \geq 3$
- $t_2 = 0$ and $t_3 = 5/612$
- large- N expansion of t_N through $1/N^9$
- deficit identity $1/4 - N^3 t_N$
- $N^3 t_N$ increases monotonically to $1/4$

The vacuum-mediated route

- $\text{ell}_N = A_N + B_N + 1/C_F$, the vacuum-mediated route at every rank

Finite-volume width

- the zone maximum of q is 12 only at even L
- $q_{\min}$ on the L -torus grid is $4 \sin^2(\pi/L)$

The orientation signs are essential

- the two band spans ARE the two incidence spectra

The second-order coefficient on a two-cube space

- connected first order vanishes by inclusion-exclusion, not by cancellation of numbers
- the connected two-cube geometry has exactly four cross-cell pairs, each -1
- the $B=6$ connected kernel is $(5/612) G_{\text{conn}} + \text{diag}$, with the certified spectrum
- the certificate's own gates pass, target-blind, with the wrong-sign control rejected
- the $B=6$ six-channel census sums to the registry's own $t_3 = 5/612$
- the $B=6$ six-channel census IS the Weingarten four-channel ledger, channel by channel
- every shared-link channel is separately proportional to the geometry, on all 56 pairs

One truncation, two constructions

- $B=6$ retains every adjacent shared-link channel; $B=4$ provably cannot
- the retention rule is $C2(\rho) + 2 C2(3) \leq B$, and both retentions it decides are equalities
- the $B=6$ six-channel census IS the Weingarten four-channel ledger, channel by channel
- FINDING: the T1 link cutoff reverses the sign of t_3 , and $14/153$ is what it omits
- FINDING: a full $T1 = B = 4$ cube Hamiltonian reproduces $-1/12$ and the reversed shell
- FINDING: the $B = 6$ cube flips the sign back to $+5/612$, closing the decisive test
- the $B = 6$ scalar misses the bridge's by exactly the same-face sextet route

What the cutoff moves, and where

- the connected diagonal is orbit-constant, and only the transporting orbit moves

The one-cube shell, and why its flat level cannot move

- the one-cube shell is $A1 + T1 + E$, and its flat level is the S^2 fundamental class
- the certificate's finite-volume fingerprints, and $29 = L^3 + 2$ is the Lean cycle count

A detached-replayed radius-two Ritz extension

- the sealed radius-two report passes its six float checks at stated tolerances
- the sealed radius-two report passes its six float checks at stated tolerances
- the sealed radius-two report passes its six float checks at stated tolerances

The unsigned-incidence comparison operator

- the C-even characteristic polynomial is $\mu(\mu - p)^2 = 4 a_1 a_2 a_3$
- $p + q = 12$: one zone function runs both sectors
- the Bloch cubic IS the finite $L = 3$ and $L = 4$ plaquette spectrum, exactly
- the C-even range $[-4, 12]$ is exact, and each edge is attained at one point only
- the C-even band touches its floor exactly on the three planes $k_j = \pi$
- the C-even band touches its floor exactly on the three planes $k_j = \pi$
- the C-even spectra at the four high-symmetry momenta

Conditional charge-even assembly

- one assembly formula gives every registered band value
- one assembly formula gives every C-even value at both orders
- the plaquette graph is 12-regular and two faces share at most one link

- both declared coincidences, checked: one is ell_N at all ranks, the other is bare

Where the rest frame is not blind

- the C-even Gamma point pins t_+ , exactly where no C-odd Gamma datum can
- the C-even curvature is $(4/3)|t_+|$ at both orders, and isotropic
- the C-even bandwidth is $16|t_+|$ at every order; the C-odd manifold width is $12|t_-|$
- at $N = 2$ the C-odd hopping vanishes and the C-even one does not

$SU(3)$ through third order

- $d_3 = 7/32 + 12 \cdot \text{leak}_3 - 4 \cdot b_3$
- leak_3 is assembled from the domino diagonal and the vacuum piece
- E_{flat} and $t(u)$ carry the ledger coefficients
- $d_- = 1/2 + 12 \cdot \text{leak}_2$, and $\text{leak}_2 = -11/306$
- FINDING: no Gamma-point datum can constrain the hopping

What the homology protects

- a boundary-factorised correction shifts the carrier without dispersing it, for generic M
- the carrier projection is where the $1/q$ comes from
- clearing the denominator reproduces the five-element numerator basis

An exact theorem for the saved historical kernel

- v10a.26 A, B, D match the sealed rationals within $2.3e-13$
- FINDING: α_{new} falls outside the corpus's own $2.3e-13$ bound

The obstruction is polynomial

- FINDING: the retained Gamma/axis data cannot identify C_{shp}
- Φ_C vanishes at Gamma along every direction
- the crosswalk is exactly scalar on the momentum axes
- on an axial cut the mixed invariants vanish and the norm divides

A second witness

- FINDING: an explicit second witness C_{alt} exhibits the C2 non-identifiability
- the C_{alt} witness IS the balanced eps-free continuation, and $\Delta_C/(A/2) = 15/32$

Where the coefficient lives

- checkpoint values at X, M, P, R
- the four extraction formulas invert the ansatz
- X is blind to B, C, D — it fixes A alone
- the shape fit's C row sums to zero, so no Gamma-anchor error can move C_{shp} at all
- a translation-local scalar shift changes nothing observable
- FINDING: C_{shp} is carried by 6 of 189 records, not spread across the kernel
- FINDING: $A = 5/48$ pins the normal sector's whole C4 contribution
- $C_{\text{normal}} = -A_{\text{normal}}/2$: the agreed axial coefficient pins the normal channel
- the two 189-record kernels agree everywhere except three amplitudes, and the on-site anchor swap moves C by exactly zero

What cannot decide it

- FINDING: the marked-cluster engine emits the Gamma scalar only — a completed 609-sweep cannot decide C_{shp}

The tier collapse, as two integers

- RETRACTED: the vertex count does NOT forbid B_{shp} and D_{shp}
- FINDING: the tier collapse is two integer cancellations at record level
- the vanishing coefficients are exactly the degree-3 ones

A near- Γ statement that does not adjudicate

- Jordan bound $q(k) \geq (4/\pi^2)|k|^2$ holds on the whole zone
- the exact zone minimum of q on $|k| \geq r$ is $4 \sin^2(r/2)$, and Jordan overstates it by $\pi/2$
- the criterion survives C2: K depends on the kernel only through $\sqrt{W_4}$
- the statement is non-vacuous only below an explicit coupling

External comparisons

- Hamer’s 1+- series matches the C-odd Gamma-point coefficients through x^3
- Hamer’s 0++ series matches the C-even Gamma-point coefficients through x^3
- the $m_n = 2^{n-1} a_n$ bridge is the $x = 2u$ conversion
- the Hamer table is pinned, and the a_4 agreement is primary-source
- the KPS 1980 string-tension table equals the certified sigma series EXACTLY
- $\sigma_n^{\text{phys}} = (-1)^n \sigma_n^{\text{raw}}$, and C5 is the $n = 3$ case
- the ratio and sigma series reproduce E_{flat} exactly
- the errata-resolved Euclidean series is doubly sourced, transcription for transcription
- FINDING: Munster’s 1985 table shifts his 1982 erratum at eighth order
- the overlap obstruction was published in 1988, and it scales
- a cross-regime paper never supplies a value

Scope

- $\Delta_L = 4 \tau(u) \sin^2(\pi/L)$ is positive and falls as L^{-2}

Channel weights in closed form

- the four channel weights follow from dimension and Casimir

The Weingarten index sums

- the shared-link weights are Weingarten, not an isotropy assumption
- $SU(3)$ Weingarten values follow from the general formula
- the fourth moment integral $|U_{11}|^4 = 1/6$ at $N = 3$

The $SU(3)$ coefficient ledger

- the manuscript’s $SU(3)$ ledger is this registry, value by value

Unsigned comparison and conditional charge-even ledger

- one assembly formula gives every C-even value at both orders
- FINDING: the certificate key ‘bandmin’ holds the band MAXIMUM, at both orders

Finite-volume chain calculation

- $d_2 d_3 = 0$ on the built complex
- $\text{rank } d_3 = L^3 - 1$ on the built complex

The two-cube geometry

- the $B=4$ comparator on the same geometry gives $-1/12$ with the reversed spectrum

References

- [1] D. A. Yarotsky, *Quasi-particles in weak perturbations of non-interacting quantum lattice systems*, [arXiv:math-ph/0411042](#) (2004).
- [2] B. Eckmann, *Harmonische Funktionen und Randwertaufgaben in einem Komplex*, *Comment. Math. Helv.* **17**, 240–255 (1944), [doi:10.1007/BF02566245](#).
- [3] J. B. Kogut and L. Susskind, *Hamiltonian formulation of Wilson’s lattice gauge theories*, *Phys. Rev. D* **11**, 395–408 (1975), [doi:10.1103/PhysRevD.11.395](#).
- [4] J. B. Kogut, D. K. Sinclair and L. Susskind, *A quantitative approach to low-energy quantum chromodynamics*, *Nucl. Phys. B* **114**, 199–236 (1976), [doi:10.1016/0550-3213\(76\)90586-1](#).
- [5] D. Weingarten, *Asymptotic behavior of group integrals in the limit of infinite rank*, *J. Math. Phys.* **19**, 999–1001 (1978), [doi:10.1063/1.523807](#).
- [6] G. Münster, *Strong coupling expansions for the mass gap in lattice gauge theories*, *Nucl. Phys. B* **190**, 439–453 (1981), with addendum **200**, 536–538 (1982) and erratum **205**, 648 (1982), [doi:10.1016/0550-3213\(81\)90570-8](#).
- [7] J. B. Kogut, R. B. Pearson and J. Shigemitsu, *The string tension, confinement and roughening in $SU(3)$ Hamiltonian lattice gauge theory*, *Phys. Lett. B* **98**, 63–68 (1981), [doi:10.1016/0370-2693\(81\)90369-5](#).
- [8] K. Seo, *Glueball mass estimate by strong coupling expansion in lattice gauge theories*, *Nucl. Phys. B* **209**, 200–216 (1982), [doi:10.1016/0550-3213\(82\)90110-9](#).
- [9] C. J. Smit, *Estimate of glueball masses from their strong coupling series in lattice QCD*, *Nucl. Phys. B* **206**, 309–320 (1982), [doi:10.1016/0550-3213\(82\)90537-5](#).
- [10] G. Münster, *Physical strong coupling expansion parameters and glueball mass ratios*, *Phys. Lett. B* **121**, 53–55 (1983), [doi:10.1016/0370-2693\(83\)90201-0](#).
- [11] G. Münster, *Effective transfer matrix for low-lying glueball states in lattice gauge theory*, *Nucl. Phys. B* **256**, 67–76 (1985), [doi:10.1016/0550-3213\(85\)90385-2](#).
- [12] B. Sutherland, *Localization of electronic wave functions due to local topology*, *Phys. Rev. B* **34**, 5208–5211 (1986), [doi:10.1103/PhysRevB.34.5208](#).
- [13] G. Schierholz, *Glueball masses in $SU(3)$ lattice gauge theory*, *Nucl. Phys. B Proc. Suppl.* **4**, 11–20 (1988), [doi:10.1016/0920-5632\(88\)90075-8](#).
- [14] C. J. Hamer, *Hamiltonian strong coupling expansions for glueball masses in $SU(3)$* , *Phys. Lett. B* **224**, 339–342 (1989), [doi:10.1016/0370-2693\(89\)91242-2](#).
- [15] F. Brandstaeter, A. S. Kronfeld and G. Schierholz, *Critical signal-to-noise ratio and glueball mass calculations*, *Nucl. Phys. B* **345**, 709–729 (1990), [doi:10.1016/0550-3213\(90\)90406-4](#).
- [16] A. Mielke, *Ferromagnetism in the Hubbard model on line graphs and further considerations*, *J. Phys. A* **24**, 3311–3321 (1991), [doi:10.1088/0305-4470/24/14/018](#).
- [17] J. Carlsson and B. H. J. McKellar, *$SU(N)$ lattice gauge theory on a single cube*, [arXiv:hep-lat/0303022](#) (2003).
- [18] B. Collins and P. Śniady, *Integration with respect to the Haar measure on unitary, orthogonal and symplectic group*, *Commun. Math. Phys.* **264**, 773–795 (2006), [doi:10.1007/s00220-006-1554-3](#).

- [19] D. L. Bergman, C. Wu and L. Balents, *Band touching from real-space topology in frustrated hopping models*, Phys. Rev. B **78**, 125104 (2008), doi:10.1103/PhysRevB.78.125104.
- [20] T. Zaslavsky, *Matrices in the theory of signed simple graphs*, in *Advances in Discrete Mathematics and Applications: Mysore, 2008*, Ramanujan Mathematical Society Lecture Notes Series **13**, 207–229 (2010), arXiv:1303.3083.
- [21] J.-W. Rhim and B.-J. Yang, *Classification of flat bands according to the band-crossing singularity of Bloch wave functions*, Phys. Rev. B **99**, 045107 (2019), doi:10.1103/PhysRevB.99.045107.
- [22] A. J. Kollár, M. Fitzpatrick, P. Sarnak and A. A. Houck, *Line-graph lattices: Euclidean and non-Euclidean flat bands, and implementations in circuit quantum electrodynamics*, Commun. Math. Phys. **376**, 1909–1956 (2020), doi:10.1007/s00220-019-03645-8.
- [23] Z. Davoudi, I. Raychowdhury and A. Shaw, *Search for efficient formulations for Hamiltonian simulation of non-Abelian lattice gauge theories*, Phys. Rev. D **104**, 074505 (2021), doi:10.1103/PhysRevD.104.074505.
- [24] S. V. Kadam, I. Raychowdhury and J. R. Stryker, *Loop-string-hadron formulation of an $SU(3)$ gauge theory with dynamical quarks*, Phys. Rev. D **107**, 094513 (2023), doi:10.1103/PhysRevD.107.094513.
- [25] K. Roychowdhury, J. Attig, S. Trebst and M. J. Lawler, *Supersymmetry on the lattice: Geometry, topology, and flat bands*, Phys. Rev. Research **6**, 043273 (2024), doi:10.1103/PhysRevResearch.6.043273.
- [26] K. T. K. Chung and M. J. P. Gingras, *2-form $U(1)$ spin liquids: A classical model and quantum aspects*, Phys. Rev. B **111**, 064417 (2025), doi:10.1103/PhysRevB.111.064417.
- [27] A. N. Ciavarella, I. M. Burbano and C. W. Bauer, *Efficient truncations of lattice gauge theory for quantum simulation*, Phys. Rev. D **112**, 054514 (2025), doi:10.1103/ylqb-phv5; arXiv:2503.11888v5.
- [28] P. Balaji, C. Conefrey-Shinozaki, P. Draper, J. K. Elhaderi, D. Gupta, L. Hidalgo and A. Lytle, *Perturbation theory, irrep truncations, and state preparation methods for quantum simulations of $SU(3)$ lattice gauge theory*, Phys. Rev. D **113**, 094505 (2026), doi:10.1103/m719-7tdf; arXiv:2509.25865v3.
- [29] HEPQIS–UIUC, *ymcirc: lattice-gauge-theory circuit and matrix-element tools*, GitHub repository, commit e9e190bfda405608de9cab71c0df0161cfcb1a10, data blob 1808849c74b36f8082e7030d074681be4bb8c0fd (accessed 28 August 2026), repository; data introduced in PR 76.
- [30] T. Hazra, *New class of exactly flat topological bands—compact localised states protected by local graph topology*, arXiv:2607.22831v1 (2026).